

Deep-Learning Image Classification with CNNs — a fully-documented model bake-off

Course: AI503 — Machine Learning | **Assignment:** W08–W09 *Extra Assignment* (Image Classification Using CNN) **Student:** Mohamed Elrashid (22002576) | **Dataset:** CIFAR-10 (one dataset, held constant)

How to read this notebook (please read me first)

This notebook assumes **no prior knowledge**. Every new word is explained the first time it appears. Each section starts with a coloured badge so you always know *what it is for*:

-  **COVERS TASK n** — this section satisfies a required part of the assignment.
-  **EXTRA** — this goes beyond the minimum (it is what turns a pass into an A).
-  **WEEK** — which teaching week the idea came from.
-  **DL Guide p.NN** — the exact page in *Deep Learning: A Comprehensive Guide* (the W05–W06 PDF) where the idea is explained.

The big idea in one sentence: the assignment asks for 3 CNNs; we instead put **~23 models on one leaderboard** and change only the architecture each time, exactly like the Week-7 `compare-all-ml-models` notebook put 18 models on one table. That is the only fair way to compare.

Assignment coverage map (read this to see nothing is missing)

Task	What the assignment asks for	Where it lives below	Status
Task 1	Describe the dataset (count, classes, size, examples, why CNN-suitable)	\$2	 Covered
Task 2	Preprocess (resize, normalize, split, augmentation) + explain why	\$3	 Covered
Task 3	Design a CNN (conv, pool, ReLU, flatten, dense, softmax)	\$4 + \$5	 Covered

Task	What the assignment asks for	Where it lives below	Status
Task 4	Train (batch size, epochs, optimizer, learning rate)	\$5 (shared harness) + \$10 (tuning)	✓ Covered
Task 5	Evaluate + compare ≥ 3 CNNs of increasing depth (acc/P/R/F1, confusion matrix, acc & loss curves)	\$5 (the 3) + \$11 (full leaderboard)	✓ Covered + exceeded (23 models)
Task 6	Discuss results — which classes are correct / confused / why	\$12	✓ Covered
Task 7	Improve (augmentation, dropout, batch-norm, transfer learning, tuning)	\$6–\$10 + \$13	✓ Covered — all five , each measured

If a cell ever fails or is skipped, its badge still tells you which task it belonged to, so you can see exactly what is and is not done.

Cell 02 — §0 · Where every idea in this notebook comes from (syllabus + PDF map)

§0 · Where every idea in this notebook comes from (syllabus + PDF map)

Two reference tables. The first maps each idea to a **teaching week**. The second maps each idea to an **exact page** of the *Deep Learning: A Comprehensive Guide* PDF (W05–W06), so any claim can be checked.

Teaching-week map

Week	Topic	Used in this notebook for
W01–W02	Logistic Regression	the idea of a linear baseline; softmax is multi-class logistic regression
W03	KNN, Decision Tree, SVM, Naive Bayes	the <i>stacking</i> meta-learner idea (a model on top of models)
W04	K-Means clustering	(not used here — images use supervised CNNs)
W05–W06	Deep Learning (CNN, RNN, LSTM)	the whole notebook — CNNs come from here
W07	Model comparison + ensembles	the single-leaderboard method this notebook copies
W08–W09	Ensemble Learning	voting & stacking in §10

DL-Guide PDF page map (file: *W05-W06/Deep Learning_ A Comprehensive Guide.pdf*, 54 pages, 12 chapters)

Idea	Chapter	Page
AI vs ML vs DL; when to use DL vs ML	Ch 1 Introduction	p.5–6
Weights, bias, activation function	Ch 2 NN Fundamentals	p.9–10
Training cycle: forward → loss → backprop → update	Ch 2	p.8
Optimizers (Adam, SGD), learning rate	Ch 2	p.13
ReLU activation	Ch 2	p.10
Feedforward / dense network (MLP)	Ch 3 FNN	p.15
Dropout · Data augmentation · Early stopping	Ch 3	p.16
Hyperparameter tuning (grid / random / Bayesian)	Ch 3	p.16
CNN full pipeline diagram (conv→pool→flatten→dense)	Ch 4 CNN	p.18
Convolution / filters / kernels / feature maps	Ch 4	p.19
Stride, padding, output-size formula	Ch 4	p.19–20
Pooling (max / average)	Ch 4	p.21
Flatten · Fully-connected · Softmax head	Ch 4	p.21
Landmark CNNs + Transfer learning (VGG, ResNet)	Ch 4	p.22
CNN applications (medical imaging, etc.)	Ch 4	p.23
Applications of DL overall	Ch 12	p.52

Ensemble voting (W08–W09 [Ensemble_Learning.pdf](#)): "majority voting" p.5, "weighted voting" p.13, why-ensembles (bias/variance) p.3. **Note:** soft-voting and stacking are standard extensions **not** shown in those slides — they are labelled as extras where used.

Honesty note on reproducibility (read once): the Week-7 notebook got bit-identical numbers from a fixed seed because scikit-learn on a CPU is deterministic. Deep learning on a **GPU is not** bit-identical by default (the GPU adds tiny numerical noise). We fix every seed we can, but expect the leaderboard *order* to be stable while the last decimal of each score may wobble by ~0.5%. The reproducible artefact here is the **ranking**, not the 4th decimal.

Cell 03 — §1 · How to run this — speed tiers (so you are never stuck waiting)

§1 · How to run this — speed tiers (so you are never stuck waiting)

Training ~23 models can take time. You control how much runs using **flags** in the next cell. A flag is just a switch: `True` = run it, `False` = skip it.

If you set...	You run	Rough time on A100	Use when
everything <code>False</code> except the ladder	the required 3+ CNNs only	~5–10 min	first sanity check
<code>RUN_REGULARIZED=True</code> , <code>RUN_MODERN=True</code>	+ dropout/BN/augment + modern blocks	~30 min	the full from-scratch story
also <code>RUN_TRANSFER=True</code>	+ VGG16/ResNet50/MobileNet/EfficientNet/ConvNeXt	+~20 min	transfer learning
also <code>RUN_FINETUNE</code> , <code>RUN_TUNING</code>	+ fine-tuning + hyper-parameter search	+~50 min	the final A-grade run

Tip for a first run: leave `QUICK_MODE = True` . It caps every model to a few epochs so the whole notebook finishes in minutes and you can confirm nothing is broken. Then set `QUICK_MODE = False` for real numbers.

You need a **GPU runtime** (Colab: *Runtime* → *Change runtime type* → *GPU*; Kaggle: *Settings* → *Accelerator* → *GPU*). On an A100/H100 turn `MIXED_PRECISION = True` for a ~2× speed-up.

Resumable. Every finished model is saved to `CKPT_DIR` . If the runtime disconnects, just run the notebook again — completed models load instantly and only the missing ones train. Set `RESUME = False` to force a clean re-run. QUICK and full runs use separate checkpoint folders, so a quick smoke-test never contaminates the real run. On Colab, point `CKPT_DIR` at a Google-Drive path (e.g. `/content/drive/MyDrive/cnn_ckpts`) so checkpoints survive a disconnect — the local disk is wiped when the runtime dies.

```
In [36]: # ===== Cell 04 – CONTROL PANEL =====
# ===== CONTROL PANEL =====
# Flip these switches to decide how much of the notebook runs.
SEED = 42 # fixed random seed (the Week-7 notebook used 42 too)

QUICK_MODE = False # True = tiny epochs for a fast smoke-test. Set False f
MIXED_PRECISION = True # True = use float16 maths on GPU (big speed-up on A100/

# --- which model groups to train (see the speed-tier table above) ---
RUN_LADDER = True # Group A: the REQUIRED depth ladder (2->3->4->5->6 conv
RUN_REGULARIZED = True # Group B: dropout / batch-norm / augmentation ablations
RUN_MODERN = True # Group C: VGG / residual / separable / global-average-p
RUN_TRANSFER = True # Group D: pretrained backbones, frozen (Task 7). Needs i
```

```

RUN_FINETUNE = True    # Group D+: unfreeze and fine-tune the best backbone (Task 4/7).
RUN_TUNING   = True    # Group E: small hyper-parameter grid search (Task 4/7).
RUN_VIT      = True    # OPTIONAL extra: Vision Transformer. Needs `pip install`

# --- resume / checkpoints ---
RESUME       = True     # True = skip models already saved in CKPT_DIR (resume a
CKPT_DIR     = 'cnn_ckpts' # where per-model results are saved. Use '/content/drive
SAVE_WEIGHTS = True     # also save each model's FULL weights (.keras) so you can

# --- sizes ---
IMG_SIZE_TL = 128       # pixels to upscale CIFAR (32px) to for the pretrained b
BATCH       = 128       # images per training step. On A100 you can raise to 256
print("Control panel set. QUICK_MODE =", QUICK_MODE, "| MIXED_PRECISION =", MIXED_P

```

Control panel set. QUICK_MODE = False | MIXED_PRECISION = True

Cell 05 — §1.1 · Setup — import the tools and make results repeatable

§1.1 · Setup — import the tools and make results repeatable

Every import below is a tool we will use. In plain words:

- **numpy** — fast number arrays (every image is just a grid of numbers).
- **pandas** — tables (our leaderboard is a pandas table).
- **matplotlib / seaborn** — drawing charts and heatmaps.
- **tensorflow.keras** — the deep-learning library. *Keras* is the friendly front-end; *TensorFlow* is the engine underneath.
- **scikit-learn metrics** — ready-made accuracy / precision / recall / F1 / confusion-matrix functions (same ones the Week-7 notebook used).

We also **fix the random seed** everywhere. Training starts from random numbers; fixing the seed means the notebook gives (almost) the same answer every run.

```

In [37]: # ===== Cell 06 — §1.1 · Setup — import the tools and make results repeatable =====
import os, gc, random, warnings, time
warnings.filterwarnings('ignore')
os.environ['TF_CPP_MIN_LOG_LEVEL'] = '2'      # hide Low-Level TensorFlow chatter
os.environ['TF_DETERMINISTIC_OPS'] = '1'     # ask the GPU to be as repeatable as

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import tensorflow as tf
from tensorflow.keras import layers, models, optimizers
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras import backend as K

from sklearn.metrics import (accuracy_score, precision_score, recall_score, f1_score,
                             roc_auc_score, confusion_matrix, classification_report

```

```

from sklearn.linear_model import LogisticRegression # used as the stacking meta-L

def set_seeds(seed=SEED):
    os.environ['PYTHONHASHSEED'] = str(seed)
    random.seed(seed); np.random.seed(seed); tf.random.set_seed(seed)

set_seeds()
sns.set_style('whitegrid')

# turn on float16 maths if asked AND a GPU is present (it does nothing useful on CP
gpus = tf.config.list_physical_devices('GPU')
if MIXED_PRECISION and gpus:
    tf.keras.mixed_precision.set_global_policy('mixed_float16')
    print('Mixed precision ON (float16).')
print('TensorFlow', tf.__version__, '| GPU available:', bool(gpus), '|', gpus)

```

Mixed precision ON (float16).

TensorFlow 2.20.0 | GPU available: True | [PhysicalDevice(name='/physical_device:GPU:0', device_type='GPU')]

Cell 07 — §1.2 · Restore checkpoints + models from Google Drive (run before training)

§1.2 · Restore checkpoints + models from Google Drive (run before training) + EXTRA

If your Colab session was wiped, this pulls any earlier **checkpoints and saved model weights** from a Google-Drive backup back to the local folder. The training cells then reload finished models instantly instead of retraining. It is safe to run every time — it only copies what is missing. The matching backup cell near the end pushes new results up to Drive. (Uses a local `CKPT_DIR` + Drive backup; for the simplest setup you can instead point `CKPT_DIR` straight at a Drive path in the control panel.)

```

In [38]: # ===== Cell 08 – RESTORE: pull checkpoints + model weights from Google Drive =====
# ===== RESTORE: pull checkpoints + model weights from Google Drive (run me before
import shutil
DRIVE_BACKUP = '/content/drive/MyDrive/cnn_backup' # where backups live on you
try:
    from google.colab import drive
    if not os.path.isdir('/content/drive/MyDrive'):
        drive.mount('/content/drive')
except Exception:
    print('(not on Colab / no Drive - skipping restore)')

if RESUME and os.path.isdir(DRIVE_BACKUP):
    n = 0
    for root, _, files in os.walk(DRIVE_BACKUP):
        for f in files:
            src = os.path.join(root, f)
            dst = os.path.join(CKPT_DIR, os.path.relpath(src, DRIVE_BACKUP))
            os.makedirs(os.path.dirname(dst), exist_ok=True)
            if not os.path.exists(dst):
                shutil.copy2(src, dst); n += 1

```

```
print(f'Restored {n} files from {DRIVE_BACKUP} -> {CKPT_DIR}/ (finished models
else:
print('No Drive backup found at', DRIVE_BACKUP, '- nothing to restore yet (it i
```

Restored 0 files from /content/drive/MyDrive/cnn_backup -> cnn_ckpts/ (finished models will load instantly).

Cell 09 — §2 · Task 1 — The dataset (CIFAR-10)

§2 · Task 1 — The dataset (CIFAR-10) COVERS TASK 1

 **WEEK** W05–W06 (Deep Learning) ·  **DL Guide** p.18, p.23 (why CNNs are made for images)

What is CIFAR-10? A famous practice dataset of **60,000 tiny colour photos**, each **32×32 pixels**, sorted into **10 everyday classes**: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck. It ships *inside* Keras, so there is nothing to download by hand.

Why is it a good fit for a CNN?  (DL Guide p.18) explains that a normal dense network would need ~150,000 inputs for even a small image and explode into millions of weights. A **CNN** instead slides small filters across the picture, so it (a) needs far fewer weights and (b) understands that nearby pixels belong together — the *spatial* structure a photo has. CIFAR-10 is the standard first dataset for exactly this lesson.

Numbers we must report for Task 1: total images, number of classes, image size, an example of each category, and the train/test split. The next two cells print all of them.

```
In [39]: # ===== Cell 10 — §2 · Task 1 — The dataset (CIFAR-10) =====
from tensorflow.keras.datasets import cifar10

# CIFAR-10 arrives already split into a 50,000-image training set and a 10,000-image
(x_train_full, y_train_full), (x_test, y_test) = cifar10.load_data()

CLASS_NAMES = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog',
               'frog', 'horse', 'ship', 'truck']
NUM_CLASSES = len(CLASS_NAMES)

print('Image shape (height, width, colour channels):', x_train_full.shape[1:])
print('Pixel value range:', x_train_full.min(), 'to', x_train_full.max(), '(0=black)')
print('Training images:', x_train_full.shape[0])
print('Test images      :', x_test.shape[0])
print('Total images     :', x_train_full.shape[0] + x_test.shape[0])
print('Number of classes:', NUM_CLASSES, '->', CLASS_NAMES)

# class balance — are the 10 classes equally common? (the Week-7 notebook checked b
unique, counts = np.unique(y_train_full, return_counts=True)
print('\nImages per class in the training set:')
for c, n in zip(unique, counts):
    print(f' {CLASS_NAMES[c]:12s}: {n}')
print('\nCIFAR-10 is perfectly balanced (5,000 of each) and has no train/test leaka
```

Image shape (height, width, colour channels): (32, 32, 3)
 Pixel value range: 0 to 255 (0=black, 255=brightest)
 Training images: 50000
 Test images : 10000
 Total images : 60000
 Number of classes: 10 -> ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']

Images per class in the training set:

airplane	: 5000
automobile	: 5000
bird	: 5000
cat	: 5000
deer	: 5000
dog	: 5000
frog	: 5000
horse	: 5000
ship	: 5000
truck	: 5000

CIFAR-10 is perfectly balanced (5,000 of each) and has no train/test leakage by construction.

```
In [40]: # ===== Cell 11 – Show one example image from every class (Task 1 asks for cat) =====
# Show one example image from every class (Task 1 asks for category examples)
fig, axes = plt.subplots(2, 5, figsize=(11, 5))
for cls, ax in enumerate(axes.flat):
    idx = np.where(y_train_full[:, 0] == cls)[0][0] # first image of this class
    ax.imshow(x_train_full[idx])
    ax.set_title(CLASS_NAMES[cls], fontsize=10)
    ax.axis('off')
plt.suptitle('CIFAR-10 – one example per class (32x32 pixels each)', y=1.02)
plt.tight_layout(); plt.show()
```

CIFAR-10 — one example per class (32x32 pixels each)



Cell 12 — §3 · Task 2 — Preprocessing

§3 · Task 2 — Preprocessing ☒ COVERS TASK 2

Four steps, each with the *why*:

1. **Normalize** — pixel values are 0–255. We divide by 255 so they become 0.0–1.0. *Why*: neural networks learn faster and more stably when inputs are small and on the same scale (same reason KNN/SVM needed scaling back in W03).
2. **Train / validation / test split** — we carve a **validation** set out of the training images. *Why*: we tune and watch the model on validation data, and only touch the test set once at the very end. Touching the test set during training would be cheating (data leakage — the same sin the Week-7 notebook warned about with duplicate rows).
3. **One-hot labels** — class 3 ("cat") becomes `[0,0,0,1,0,0,0,0,0,0]`. *Why*: the softmax output layer produces 10 probabilities, so the true answer must also be 10 numbers.
4. **Data augmentation** — randomly flip/rotate/zoom images *during training only*. (p.16) *Why*: it shows the model slightly different versions of each photo every epoch, so it cannot simply memorise the training set (it fights overfitting). We will also **measure** how much augmentation helps by turning it into its own leaderboard row in §6.

Two image scales. From-scratch CNNs train on the native **32×32** images in `[0,1]`. The pretrained backbones in §8 were built for big ImageNet photos, so they get the **raw 0–255** images upscaled to `IMG_SIZE_TL`, and each backbone applies *its own* normalisation internally. Mixing those up quietly wrecks accuracy, so we keep both versions of the data.

```
In [41]: # ===== Cell 13 – §3 · Task 2 – Preprocessing =====
from sklearn.model_selection import train_test_split

# ---- version A: scaled [0,1] floats for the FROM-SCRATCH CNNs ----
x_train_full_s = x_train_full.astype('float32') / 255.0
x_test_s       = x_test.astype('float32') / 255.0

# ---- version B: raw [0,255] floats for the PRETRAINED backbones (they scale inter
x_train_full_r = x_train_full.astype('float32')
x_test_r       = x_test.astype('float32')

# ---- one-hot labels ----
y_train_full_oh = tf.keras.utils.to_categorical(y_train_full, NUM_CLASSES)
y_test_oh       = tf.keras.utils.to_categorical(y_test, NUM_CLASSES)

# ---- carve a validation set out of the 50k training images (stratified = keep cla
idx_tr, idx_va = train_test_split(np.arange(len(x_train_full_s)), test_size=5000,
                                random_state=SEED, stratify=y_train_full[:, 0])

Xtr_s, Xva_s = x_train_full_s[idx_tr], x_train_full_s[idx_va] # scaled, for from-
Xtr_r, Xva_r = x_train_full_r[idx_tr], x_train_full_r[idx_va] # raw, for transfer
ytr_oh, yva_oh = y_train_full_oh[idx_tr], y_train_full_oh[idx_va]
ytr_int = y_train_full[idx_tr, 0] # plain integer la
```

```
# integer truth vectors used by the metrics later
y_test_int = y_test[:, 0]
y_val_int = y_train_full[idx_va, 0]

print('Train images      : ', Xtr_s.shape[0])
print('Validation images: ', Xva_s.shape[0])
print('Test images       : ', x_test_s.shape[0])
print('Label shape (one-hot): ', ytr_oh.shape, '-> 10 numbers per image')
```

```
Train images      : 45000
Validation images: 5000
Test images       : 10000
Label shape (one-hot): (45000, 10) -> 10 numbers per image
```

```
In [42]: # ===== Cell 14 – Build the augmentation pipeline as Keras Layers. These Layer =====
# Build the augmentation pipeline as Keras Layers. These layers ONLY change images
data_augment = models.Sequential([
    layers.RandomFlip('horizontal'), # a mirrored cat is still a cat
    layers.RandomRotation(0.08),    # small tilt
    layers.RandomZoom(0.10),        # slight zoom in/out
    layers.RandomTranslation(0.10, 0.10) # small shift
], name='data_augment')

# Show the SAME image augmented 5 different ways, so you can see what the model see
sample = Xtr_s[0:1]
fig, axes = plt.subplots(1, 6, figsize=(13, 2.6))
axes[0].imshow(sample[0]); axes[0].set_title('original'); axes[0].axis('off')
for i in range(1, 6):
    aug = data_augment(sample, training=True)[0].numpy().astype('float32')
    axes[i].imshow(np.clip(aug, 0, 1)); axes[i].set_title(f'augmented {i}'); axes[i].axis('off')
plt.suptitle('Data augmentation – 5 random versions of one training image (DL Guide p.16)')
plt.tight_layout(); plt.show()
```



Cell 15 — §3.1 · The shared scoreboard and training helper (the heart of the notebook)

§3.1 · The shared scoreboard and training helper (the heart of the notebook)

Just like the Week-7 notebook had **one** `evaluate()` function and **one** results dictionary that every model wrote into, we build the same thing here — but upgraded from 2-class to **10-class** images. That upgrade matters and is the single most important code in the notebook:

- `average='macro'` on precision/recall/F1 — score each of the 10 classes, then average. (The Week-7 version was binary; 10 classes needs macro.)
- `roc_auc_score(..., multi_class='ovr')` — "one-vs-rest" ROC for 10 classes.
- We read **train accuracy, validation accuracy, train loss and validation loss** straight from the training `history` — Task 5 asks for all four as reported numbers, not just curves.
- We compute **test** accuracy/precision/recall/F1 on the untouched 10k test set.

Every model therefore lands one tidy row on the same scoreboard, and we can sort them all fairly at the end.

```
In [43]: # ===== Cell 16 - §3.1 · The shared scoreboard and training helper (the heart of the
results = {} # name -> dict of scores + predictions (our Leaderboard rows)
histories = {} # name -> per-epoch training history (for the accuracy/loss curves)

# ---- resume support: save each finished model to disk; reload it instead of retrain
import pickle
_CKPT = os.path.join(CKPT_DIR, 'quick' if QUICK_MODE else 'full') # keep QUICK and
os.makedirs(_CKPT, exist_ok=True)
def _ckpt_path(name):
    safe = name.replace('/', '_').replace(' ', '_').replace(':', '')
    return os.path.join(_CKPT, safe + '.pkl')
def _save_ckpt(name):
    try:
        with open(_ckpt_path(name), 'wb') as f:
            pickle.dump({'row': results[name], 'hist': histories[name]}, f)
    except Exception as e:
        print(' (checkpoint save failed:', e, ')')
def _load_ckpt(name):
    p = _ckpt_path(name)
    if RESUME and os.path.exists(p):
        with open(p, 'rb') as f:
            d = pickle.load(f)
            results[name] = d['row']; histories[name] = d['hist']
        return True
    return False

# ---- optional: also save FULL model weights (.keras) so a trained model can be reloaded
_WCKPT = os.path.join(_CKPT, 'weights')
os.makedirs(_WCKPT, exist_ok=True)
def _weights_path(name):
    safe = name.replace('/', '_').replace(' ', '_').replace(':', '')
    return os.path.join(_WCKPT, safe + '.keras')
def _reload_weights(name):
    p = _weights_path(name)
    if os.path.exists(p):
        try:
            return tf.keras.models.load_model(p, safe_mode=False, compile=False)
        except Exception as e:
            print(' (could not reload weights for', name, '->', str(e)[:80], ')')
    return None

def make_callbacks(patience=6):
```

```

# EarlyStopping = stop when validation stops improving (and keep the best weigh
# ReduceLROnPlateau = lower the learning rate when progress stalls (DL Guide p.
return [EarlyStopping(monitor='val_loss', patience=patience, restore_best_weigh
        ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=3, min_lr=1e

def epochs_for(default):
    return 3 if QUICK_MODE else default # QUICK_MODE shrinks every training run

def run_model(name, group, build_fn, use_raw=False, augment=False,
              epochs=30, lr=1e-3, keep_model=False, fine_tune=None):
    '''Train one model, score it on val+test, and store one leaderboard row.
    use_raw : feed raw [0,255] images (for pretrained backbones) instead of [0
    augment : prepend the augmentation layers (training only).
    fine_tune : optional dict {"unfreeze": base_model, "epochs": n, "lr": small}
    Returns the trained model (kept only if keep_model=True, else freed to save

    if _load_ckpt(name):
        r = results[name]
        print(f'{name:34s} [resumed from checkpoint] test_acc={r["test_acc"]:.3f}
        return _reload_weights(name) if keep_model else None
    K.clear_session(); set_seeds(); t0 = time.time()
    Xtr, Xva, Xte = (Xtr_r, Xva_r, x_test_r) if use_raw else (Xtr_s, Xva_s, x_test_

    net = build_fn()
    if augment:
        _inp = layers.Input(Xtr.shape[1:])
        net = models.Model(_inp, net(data_augment(_inp)), name='augmented')
    net.compile(optimizer=optimizers.Adam(lr), loss='categorical_crossentropy', met
    h = net.fit(Xtr, ytr_oh, validation_data=(Xva, yva_oh),
               epochs=epochs_for(epochs), batch_size=BATCH,
               callbacks=make_callbacks(), verbose=0)
    hist = {k: list(v) for k, v in h.history.items()}

    # optional fine-tuning phase: unfreeze a pretrained base and train gently with
    if fine_tune is not None and not QUICK_MODE:
        for _lyr in net.layers: # unfreeze everything, including the nest
            _lyr.trainable = True
        net.compile(optimizer=optimizers.Adam(fine_tune.get('lr', 1e-5)),
                   loss='categorical_crossentropy', metrics=['accuracy'])
        h2 = net.fit(Xtr, ytr_oh, validation_data=(Xva, yva_oh),
                   epochs=epochs_for(fine_tune.get('epochs', 8)), batch_size=BATC
                   callbacks=make_callbacks(), verbose=0)
        for k, v in h2.history.items():
            hist.setdefault(k, []).extend(list(v))

    _score_and_store(name, group, net, hist, Xva, Xte, time.time() - t0)
    if keep_model:
        return net
    del net; gc.collect(); K.clear_session()
    return None

def _score_and_store(name, group, net, hist, Xva, Xte, secs):
    prob_test = net.predict(Xte, batch_size=256, verbose=0) # shape (10000, 10)
    prob_val = net.predict(Xva, batch_size=256, verbose=0)
    pred_test = prob_test.argmax(1)

    row = {

```

```

    'group'      : group,
    'train_acc'  : hist['accuracy'][-1],
    'val_acc'    : hist['val_accuracy'][-1],
    'test_acc'   : accuracy_score(y_test_int, pred_test),
    'train_loss' : hist['loss'][-1],
    'val_loss'   : hist['val_loss'][-1],
    'precision'  : precision_score(y_test_int, pred_test, average='macro'),
    'recall'     : recall_score(y_test_int, pred_test, average='macro'),
    'f1'         : f1_score(y_test_int, pred_test, average='macro'),
    'roc_auc'    : roc_auc_score(y_test_oh, prob_test, multi_class='ovr', average='macro'),
    'seconds'    : round(secs, 1),
    'y_pred'     : pred_test,
    'prob_test'  : prob_test.astype('float32'),
    'prob_val'   : prob_val.astype('float32'),
}
results[name] = row; histories[name] = hist
_save_ckpt(name)
if SAVE_WEIGHTS:
    try:
        net.save(_weights_path(name))
    except Exception as e:
        print(' (weight save failed:', e, ')')
print(f'{name:34s} test_acc={row["test_acc"]:.3f}  F1={row["f1"]:.3f}  '
      f'AUC={row["roc_auc"]:.3f}  ({row["seconds"]}s)')

print('Scoreboard ready. Each model will print one line and store one row.')

```

Scoreboard ready. Each model will print one line and store one row.

Cell 17 — §4 · Task 3 — The CNN building blocks, explained from zero

§4 · Task 3 — The CNN building blocks, explained from zero COVERS TASK 3

 WEEK W05–W06 ·  DL Guide Ch 4, p.18–21

Before we build models, here is every Lego brick a CNN is made of. The required pieces (conv, pool, ReLU, flatten, dense, softmax) are all here.

Brick	In plain words	 Page
Convolution (Conv2D)	A small window (a <i>filter</i> / <i>kernel</i> , e.g. 3×3) slides over the image and looks for one pattern (an edge, a colour blob). Each filter makes one feature map .	p.19
Filter / kernel	The little grid of numbers that <i>is</i> the pattern detector. A layer has many filters, so it finds many patterns at once.	p.19
Feature map	The output image showing <i>where</i> a filter fired.	p.19
ReLU	After each conv, replace negatives with 0: <code>max(0, x)</code> . Adds non-linearity so the network can learn complex shapes.	p.10

Brick	In plain words	Page
Pooling (MaxPooling2D)	Shrink the feature map by keeping the strongest value in each 2×2 window. Makes the model smaller and a little position-proof.	p.21
Flatten	Unroll the final 2-D feature maps into one long 1-D list of numbers, so a normal dense layer can read it.	p.21
Dense (fully connected)	A classic neuron layer that mixes all features to make the decision.	p.21
Softmax	The final layer turns raw scores into 10 probabilities that add up to 1. The biggest one is the predicted class.	p.21

The standard CNN recipe (p.18): Input → [Conv → ReLU → Pool] × N → Flatten → Dense → Softmax . Every model in Group A is exactly this recipe with a different number of [Conv→ReLU→Pool] blocks.

Why the last layer says dtype='float32' : when mixed precision is ON, the network does its maths in fast float16. But probabilities need full float32 precision to stay accurate, so we force *only* the final softmax back to float32. You will see this on every model.

```
In [44]: # ===== Cell 18 – The parametric from-scratch CNN factory used for the whole d =====
# The parametric from-scratch CNN factory used for the whole depth ladder AND the r
def build_cnn(n_blocks=3, base_filters=32, dropout=0.0, batchnorm=False, name=None)
    # n_blocks = how many [Conv -> (BN) -> Pool] blocks. This is the "depth" we inc
    m = models.Sequential(name=name or f'cnn_{n_blocks}blocks')
    m.add(layers.Input((32, 32, 3)))
    f = base_filters
    pools_used = 0
    for i in range(n_blocks):
        m.add(layers.Conv2D(f, 3, padding='same', activation='relu')) # Conv + Re
        if batchnorm:
            m.add(layers.BatchNormalization()) # steadies
        if pools_used < 4: # at most 4 pools: 32->16->8->4->2 (nev
            m.add(layers.MaxPooling2D()) # pooling
            pools_used += 1
        f = min(f * 2, 256) # double the filters each block (common
    m.add(layers.Flatten()) # flatten (
    m.add(layers.Dense(128, activation='relu')) # dense head
    if dropout > 0:
        m.add(layers.Dropout(dropout)) # dropout (D
    m.add(layers.Dense(NUM_CLASSES, activation='softmax', dtype='float32')) # soft
    return m

# print the Layer table of a 4-block CNN so you can SEE the recipe
build_cnn(4).summary()
```

Model: "cnn_4blocks"

Layer (type)	Output Shape	Par
conv2d_8 (Conv2D)	(None, 32, 32, 32)	
max_pooling2d_8 (MaxPooling2D)	(None, 16, 16, 32)	
conv2d_9 (Conv2D)	(None, 16, 16, 64)	18
max_pooling2d_9 (MaxPooling2D)	(None, 8, 8, 64)	
conv2d_10 (Conv2D)	(None, 8, 8, 128)	73
max_pooling2d_10 (MaxPooling2D)	(None, 4, 4, 128)	
conv2d_11 (Conv2D)	(None, 4, 4, 256)	295
max_pooling2d_11 (MaxPooling2D)	(None, 2, 2, 256)	
flatten_2 (Flatten)	(None, 1024)	
dense_4 (Dense)	(None, 128)	131
dense_5 (Dense)	(None, 10)	1

Total params: 520,906 (1.99 MB)

Trainable params: 520,906 (1.99 MB)

Non-trainable params: 0 (0.00 B)

Cell 19 — §5 · Group A — the required depth ladder (Task 5 core)

§5 · Group A — the required depth ladder (Task 5 core) COVERS TASK 5 & TASK 3 & TASK 4

 WEEK W05–W06 ·  DL Guide p.18

The assignment requires **three** CNNs of growing depth: Model 1 = 2 conv layers, Model 2 = 3–4, Model 3 = 5+. We do that **and** add the in-between steps, training **2 → 3 → 4 → 5 → 6** conv layers. This is the "add one layer, re-train, measure" idea — and it produces one clean *accuracy-vs-depth* chart.


Honest expectation (write this down before running): accuracy should **rise** from 2 to ~4 layers as the network learns richer features, then **plateau or dip** at 5–6 layers because (a) after a few poolings a 32×32 image has almost no spatial size left, (b) deep plain networks suffer vanishing gradients, and (c) more layers = more parameters = more overfitting on limited data. That plateau is the whole lesson — and the reason Groups B–D exist.

Training settings (Task 4): optimizer = **Adam**, learning rate = **0.001**, batch size = **128**, loss = **categorical cross-entropy**, with **early stopping** so each model trains just long enough.

```
In [45]: # ===== Cell 20 – Train the depth ladder. We keep the 4-conv model in memory f =====
# Train the depth ladder. We keep the 4-conv model in memory for the Grad-CAM picture
KEEP_FOR_GRADCAM = None
LADDER = []
if RUN_LADDER:
    for depth in [2, 3, 4, 5, 6]:
        nm = f'CNN-{depth}conv'
        keep = (depth == 4)
        m = run_model(nm, 'A_ladder', (lambda d=depth: build_cnn(d)),
                      epochs=30, lr=1e-3, keep_model=keep)
        if keep:
            KEEP_FOR_GRADCAM = m
            LADDER.append((depth, nm))
else:
    print('RUN_LADDER is False – skipping the required ladder (set it True).')
```


CNN-2conv 06	[resumed from checkpoint]	test_acc=0.704	F1=0.7
CNN-3conv 30	[resumed from checkpoint]	test_acc=0.732	F1=0.7
CNN-4conv 32	[resumed from checkpoint]	test_acc=0.735	F1=0.7
CNN-5conv 95	[resumed from checkpoint]	test_acc=0.700	F1=0.6
CNN-6conv 05	[resumed from checkpoint]	test_acc=0.704	F1=0.7

```
In [46]: # ===== Cell 21 – Chart 1: accuracy vs depth (the headline picture for Task 5' =====
# Chart 1: accuracy vs depth (the headline picture for Task 5's architecture comparison)
if LADDER:
    depths = [d for d, _ in LADDER]
    accs = [results[n]['test_acc'] for _, n in LADDER]
    plt.figure(figsize=(8, 5))
    plt.plot(depths, accs, 'o-', color='#4472C4', linewidth=2, markersize=9)
    for d, a in zip(depths, accs):
        plt.text(d, a + 0.004, f'{a:.3f}', ha='center', fontsize=10, fontweight='bold')
    plt.xlabel('Number of convolutional layers (depth)')
    plt.ylabel('Test accuracy')
    plt.title('Accuracy vs depth – rises, then plateaus (the core Task-5 finding)')
    plt.xticks(depths); plt.tight_layout(); plt.show()
```

```
In [47]: # ===== Cell 22 – Chart 2: train & validation accuracy and Loss curves for the =====
# Chart 2: train & validation accuracy and loss curves for the three REQUIRED model
required = [('Model 1 = 2 conv', 'CNN-2conv'),
            ('Model 2 = 4 conv', 'CNN-4conv'),
            ('Model 3 = 6 conv', 'CNN-6conv')]
colors = ['#4472C4', '#ED7D31', '#70AD47']
fig, ax = plt.subplots(1, 2, figsize=(13, 4.6))
for (label, nm), c in zip(required, colors):
    if nm in histories and histories[nm]:
        h = histories[nm]
        ax[0].plot(h['accuracy'], color=c, label=f'{label} (train)')
        ax[0].plot(h['val_accuracy'], color=c, linestyle='--', label=f'{label} (val)')
        ax[1].plot(h['loss'], color=c, label=f'{label} (train)')
        ax[1].plot(h['val_loss'], color=c, linestyle='--', label=f'{label} (val)')
ax[0].set_title('Accuracy curves (solid=train, dashed=val)'); ax[0].set_xlabel('epoch')
ax[1].set_title('Loss curves (solid=train, dashed=val)'); ax[1].set_xlabel('epoch')
plt.suptitle('Required 3 CNNs – training vs validation (a widening gap = overfitting)')
plt.tight_layout(); plt.show()
```

Required 3 CNNs — training vs validation (a widening gap = overfitting)



§6 · Group B — regularisation ablations (Task 7, part 1) COVERS TASK 7 EXTRA

 **WEEK** W05–W06 ·  **DL Guide** p.16 (dropout, augmentation)

Task 7 says "improve using one or more of: augmentation, dropout, batch-norm...". Instead of picking one, we add **each one separately** to the same 5-conv base, so the leaderboard shows exactly **how much each trick is worth**.

- **Dropout**  (p.16): randomly switch off neurons during training so the model cannot lean on any single one → fights overfitting.
- **Batch normalisation**: re-centre each layer's outputs so training is faster and steadier. *(Standard technique; not in the W05–W06 slides — labelled extra.)*
- **Augmentation**  (p.16): the random flips/rotations from §3, now switched on.
- **All three together**: the fully-regularised custom net.

```
In [48]: # ===== Cell 24 — §6 · Group B — regularisation ablations (Task 7, part 1) =====
if RUN_REGULARIZED:
    run_model('Deep+Dropout', 'B_reg', lambda: build_cnn(5, dropout=0.4))
    run_model('Deep+BatchNorm', 'B_reg', lambda: build_cnn(5, batchnorm=True))
    run_model('Deep+Augment', 'B_reg', lambda: build_cnn(5, augment=True))
    run_model('Deep+All(BN+Drop+Aug)', 'B_reg', lambda: build_cnn(5, dropout=0.4, ba
else:
    print('RUN_REGULARIZED is False — skipping Group B.')
```

```
Deep+Dropout          [resumed from checkpoint]  test_acc=0.741  F1=0.7
42
Deep+BatchNorm        [resumed from checkpoint]  test_acc=0.707  F1=0.7
07
Deep+Augment          [resumed from checkpoint]  test_acc=0.770  F1=0.7
66
Deep+All(BN+Drop+Aug) [resumed from checkpoint]  test_acc=0.801  F1=0.8
01
```

§7 · Group C — modern CNN blocks, built from scratch EXTRA

 **WEEK** W05–W06 ·  **DL Guide** p.22 (landmark architectures: VGG, ResNet)

These four show *design ideas* from famous networks, built small enough to train from scratch:

- **VGG-style**: two conv layers before each pool (deeper, richer blocks).  p.22

- **Residual / skip connection** (the ResNet idea  p.22): add the block's input to its output, so gradients have a shortcut and very deep nets still train.
- **Depthwise-separable conv** (the MobileNet idea): a cheaper convolution that splits the work, far fewer parameters.
- **All-convolutional + Global Average Pooling**: no flatten, no big dense layer — average each final feature map straight into a class score. Tiny and overfit-resistant.

```
In [49]: # ===== Cell 26 - §7 · Group C - modern CNN blocks, built from scratch =====
def build_vgg():
    m = models.Sequential(name='vgg_style')
    m.add(layers.Input((32, 32, 3)))
    for f in [32, 64, 128]:
        m.add(layers.Conv2D(f, 3, padding='same', activation='relu'))
        m.add(layers.Conv2D(f, 3, padding='same', activation='relu')) # two convs
        m.add(layers.BatchNormalization())
        m.add(layers.MaxPooling2D())
    m.add(layers.Flatten())
    m.add(layers.Dense(128, activation='relu'))
    m.add(layers.Dropout(0.4))
    m.add(layers.Dense(NUM_CLASSES, activation='softmax', dtype='float32'))
    return m

def build_residual():
    inp = layers.Input((32, 32, 3))
    x = layers.Conv2D(32, 3, padding='same', activation='relu')(inp)
    x = layers.BatchNormalization()(x)
    def res_block(t, f):
        shortcut = t
        y = layers.Conv2D(f, 3, padding='same', activation='relu')(t)
        y = layers.BatchNormalization()(y)
        y = layers.Conv2D(f, 3, padding='same')(y)
        y = layers.BatchNormalization()(y)
        if shortcut.shape[-1] != f:
            shortcut = layers.Conv2D(f, 1, padding='same')(shortcut) # match ch
        out = layers.Activation('relu')(layers.add([shortcut, y])) # the skip
        return layers.MaxPooling2D()(out)
    x = res_block(x, 64)
    x = res_block(x, 128)
    x = layers.GlobalAveragePooling2D()(x)
    x = layers.Dense(128, activation='relu')(x)
    out = layers.Dense(NUM_CLASSES, activation='softmax', dtype='float32')(x)
    return models.Model(inp, out, name='residual_mini')

def build_separable():
    m = models.Sequential(name='separable')
    m.add(layers.Input((32, 32, 3)))
    m.add(layers.Conv2D(32, 3, padding='same', activation='relu'))
    for f in [64, 128, 256]:
        m.add(layers.SeparableConv2D(f, 3, padding='same', activation='relu'))
        m.add(layers.BatchNormalization())
        m.add(layers.MaxPooling2D())
    m.add(layers.GlobalAveragePooling2D())
    m.add(layers.Dense(NUM_CLASSES, activation='softmax', dtype='float32'))
    return m
```

```

def build_allconv_gap():
    m = models.Sequential(name='allconv_gap')
    m.add(layers.Input((32, 32, 3)))
    for f in [48, 96, 192]:
        m.add(layers.Conv2D(f, 3, padding='same', activation='relu'))
        m.add(layers.Conv2D(f, 3, strides=2, padding='same', activation='relu')) #
        m.add(layers.BatchNormalization())
    m.add(layers.Conv2D(NUM_CLASSES, 1, padding='same'))
    m.add(layers.GlobalAveragePooling2D())
    m.add(layers.Activation('softmax', dtype='float32'))
    return m

if RUN_MODERN:
    gm = run_model('VGG-style (2xconv blocks)', 'C_modern', build_vgg, keep_model=True)
    if KEEP_FOR_GRADCAM is None:
        KEEP_FOR_GRADCAM = gm # fallback Grad-CAM model
    run_model('Residual-mini (skip conn.)', 'C_modern', build_residual)
    run_model('Depthwise-separable', 'C_modern', build_separable)
    run_model('All-conv + GAP', 'C_modern', build_allconv_gap)
else:
    print('RUN_MODERN is False - skipping Group C.')

```

VGG-style (2xconv blocks)	[resumed from checkpoint]	test_acc=0.765	F1=0.766
Residual-mini (skip conn.)	[resumed from checkpoint]	test_acc=0.754	F1=0.755
Depthwise-separable	[resumed from checkpoint]	test_acc=0.732	F1=0.727
All-conv + GAP	[resumed from checkpoint]	test_acc=0.771	F1=0.772

Cell 27 — §8 · Group D — transfer learning with pretrained backbones (Task 7, part 2)

§8 · Group D — transfer learning with pretrained backbones (Task 7, part 2) COVERS TASK 7 EXTRA

 WEEK W05–W06 ·  DL Guide p.22 (VGG, ResNet)

Transfer learning means: take a giant network already trained on millions of ImageNet photos, and reuse its learned "eyes". We **freeze** its body (keep its weights) and only train a new 10-class head on top.  (p.22) names VGGNet and ResNet as exactly these reusable networks.

Two phases:

1. **Frozen** — train only the new head. Fast, and usually already beats from-scratch CNNs.
2. **Fine-tuning** (`RUN_FINETUNE`) — then *unfreeze* the body and train gently with a tiny learning rate (1e-5) so it adapts to CIFAR without forgetting. This usually gives the best score on the board.

The pretrained images were big, so we upscale CIFAR from 32×32 to `IMG_SIZE_TL`. Each backbone also applies **its own** normalisation (that is why VGG/ResNet/MobileNet get a `preprocess_input`, while EfficientNet/ConvNeXt already include it). Needs internet the first time to download the weights.

```
In [50]: # ===== Cell 28 – Load whichever backbones this TensorFlow version provides (s =====
# Load whichever backbones this TensorFlow version provides (skip any that are miss
import importlib
BACKBONES = {}
_specs = [('VGG16', 'vgg16', 'VGG16', False),
          ('ResNet50', 'resnet50', 'ResNet50', False),
          ('MobileNetV2', 'mobilenet_v2', 'MobileNetV2', False),
          ('EfficientNetV2S', 'efficientnet_v2', 'EfficientNetV2S', True),
          ('ConvNeXtTiny', 'convnext', 'ConvNeXtTiny', True)]
for disp, mod, ctor, builtin_prep in _specs:
    try:
        m = importlib.import_module(f'tensorflow.keras.applications.{mod}')
        BACKBONES[disp] = (getattr(m, ctor), None if builtin_prep else getattr(m, '
    except Exception as e:
        print(f' (skip {disp}: {e})')
print('Backbones available:', list(BACKBONES.keys()))

def build_transfer(backbone_ctor, preprocess_fn, size=IMG_SIZE_TL):
    base = backbone_ctor(include_top=False, weights='imagenet',
                        input_shape=(size, size, 3), pooling='avg')
    base.trainable = False # FREEZE the body
    inp = layers.Input((32, 32, 3)) # raw 0-255 CIFAR image
    x = layers.Resizing(size, size)(inp) # upscale to the size t
    if preprocess_fn is not None:
        x = layers.Lambda(lambda t: preprocess_fn(tf.cast(t, 'float32')))(x) # bac
    x = base(x, training=False)
    x = layers.Dropout(0.3)(x)
    out = layers.Dense(NUM_CLASSES, activation='softmax', dtype='float32')(x) # ne
    return models.Model(inp, out, name=f'{backbone_ctor.__name__}_TL')

if RUN_TRANSFER:
    for bname, (ctor, prep) in BACKBONES.items():
        run_model(f'{bname} (frozen TL)', 'D_transfer',
                (lambda c=ctor, p=prep: build_transfer(c, p)),
                use_raw=True, epochs=12, lr=1e-3)
else:
    print('RUN_TRANSFER is False – skipping pretrained backbones.')
```

```

Backbones available: ['VGG16', 'ResNet50', 'MobileNetV2', 'EfficientNetV2S', 'ConvNe
XtTiny']
VGG16 (frozen TL) [resumed from checkpoint] test_acc=0.842 F1=0.8
40
ResNet50 (frozen TL) [resumed from checkpoint] test_acc=0.889 F1=0.8
89
MobileNetV2 (frozen TL) [resumed from checkpoint] test_acc=0.858 F1=0.8
58
EfficientNetV2S (frozen TL) [resumed from checkpoint] test_acc=0.875 F1=0.8
75
ConvNeXtTiny (frozen TL) [resumed from checkpoint] test_acc=0.915 F1=0.9
15

```

```

In [51]: # ===== Cell 29 - Fine-tune ONE champion backbone (ResNet50): phase-1 frozen h =====
# Fine-tune ONE champion backbone (ResNet50): phase-1 frozen head, then unfreeze an
if RUN_FINETUNE and 'ResNet50' in BACKBONES:
    ctor, prep = BACKBONES['ResNet50']
    run_model('ResNet50 (fine-tuned)', 'D_transfer',
              (lambda: build_transfer(ctor, prep)),
              use_raw=True, epochs=12, lr=1e-3,
              fine_tune={'epochs': 8, 'lr': 1e-5})
else:
    print('RUN_FINETUNE is False (or ResNet50 unavailable) - skipping fine-tuning.'

```

```

ResNet50 (fine-tuned) [resumed from checkpoint] test_acc=0.922 F1=0.9
22

```

Cell 30 — §9 · Group E — hyper-parameter tuning (Task 4 / Task 7)

§9 · Group E — hyper-parameter tuning (Task 4 / Task 7) COVERS TASK 7 EXTRA

 **WEEK** W05–W06 ·  **DL Guide p.16** (grid / random / Bayesian search)

The Week-7 notebook used `GridSearchCV` to tune classic models. The CNN version of that is a small **grid search** over a couple of settings — learning rate and dropout — on the 4-conv CNN.  (p.16) lists grid / random / Bayesian as the three search strategies; for a serious run, the `keras-tuner` library does this automatically, but a hand grid keeps the notebook dependency-free.

```

In [52]: # ===== Cell 31 - §9 · Group E - hyper-parameter tuning (Task 4 / Task 7) =====
if RUN_TUNING:
    for lr_ in [1e-3, 5e-4]:
        for dr_ in [0.3, 0.5]:
            run_model(f'Grid lr={lr_: .0e} drop={dr_}', 'E_tuning',
                      (lambda d=dr_: build_cnn(4, dropout=d)), epochs=25, lr=lr_)
else:
    print('RUN_TUNING is False - skipping the hyper-parameter grid.')

```

Grid lr=1e-03 drop=0.3 47	[resumed from checkpoint]	test_acc=0.748	F1=0.7
Grid lr=1e-03 drop=0.5 48	[resumed from checkpoint]	test_acc=0.746	F1=0.7
Grid lr=5e-04 drop=0.3 29	[resumed from checkpoint]	test_acc=0.733	F1=0.7
Grid lr=5e-04 drop=0.5 53	[resumed from checkpoint]	test_acc=0.754	F1=0.7

Cell 32 — §10 · Group F — ensembles: combine the best models

§10 · Group F — ensembles: combine the best models + EXTRA (Task 7)

🎓 **WEEK** W08–W09 (Ensemble Learning) · 📄 `Ensemble_Learning.pdf` p.3 (why ensembles help), p.5 ("majority voting"), p.13 ("weighted voting")

An **ensemble** combines several models so their mistakes cancel out 📄 (p.3). We take the **top-3 models so far** and combine them three ways:

- **Soft vote** — average the 10 probabilities from each model, then pick the biggest. (*Averaging is the regression cousin of the "majority voting" on p.5.*)
- **Hard vote** — each model votes for one class; the majority wins 📄 (p.5).
- **Stacking** — train a small Logistic Regression (a W01–W02 / W03 model!) **on top of** the base models' outputs. (*Stacking is a standard extension, not shown in the W08–W09 slides — labelled extra.*)

No cheating: the stacking meta-learner is trained on the **validation** predictions and judged on the **untouched test** set, so it never sees the test answers in advance — the same data-hygiene rule the Week-7 notebook followed with `StackingClassifier(cv=5)`.

```
In [53]: # ===== Cell 33 — §10 · Group F — ensembles: combine the best models =====
def _store_pred(name, group, prob_test=None, pred=None):
    # store an ensemble row (no training history, so train/val columns are blank)
    if pred is None:
        pred = prob_test.argmax(1)
    row = {'group': group, 'train_acc': np.nan, 'val_acc': np.nan,
          'test_acc': accuracy_score(y_test_int, pred),
          'train_loss': np.nan, 'val_loss': np.nan,
          'precision': precision_score(y_test_int, pred, average='macro'),
          'recall': recall_score(y_test_int, pred, average='macro'),
          'f1': f1_score(y_test_int, pred, average='macro'),
          'roc_auc': (roc_auc_score(y_test_oh, prob_test, multi_class='ovr', average='weighted'
                                   if prob_test is not None else np.nan),
                     'seconds': 0.0, 'y_pred': pred,
                     'prob_test': None if prob_test is None else prob_test.astype('float32'),
                     'prob_val': None}
    results[name] = row; histories[name] = {}
    print(f'{name:34s} test_acc={row["test_acc"]:.3f} F1={row["f1"]:.3f}')
```

```

trained = [n for n, r in results.items() if r.get('prob_val') is not None]
if len(trained) >= 2:
    top = sorted(trained, key=lambda n: results[n]['f1'], reverse=True)[:3]
    print('Combining top-3 models:', top, '\n')

    soft = np.mean([results[n]['prob_test'] for n in top], axis=0) # soft
    _store_pred('Ensemble: Soft-Vote (top3)', 'F_ensemble', prob_test=soft)

    stack_preds = np.stack([results[n]['y_pred'] for n in top], axis=1) # hard
    hard = np.array([np.bincount(r, minlength=NUM_CLASSES).argmax() for r in stack_preds])
    _store_pred('Ensemble: Hard-Vote (top3)', 'F_ensemble', pred=hard)

    Xv = np.concatenate([results[n]['prob_val'] for n in top], axis=1) # stack
    Xt = np.concatenate([results[n]['prob_test'] for n in top], axis=1)
    meta = LogisticRegression(max_iter=2000)
    meta.fit(Xv, y_val_int)
    _store_pred('Ensemble: Stacking (LogReg)', 'F_ensemble',
                prob_test=meta.predict_proba(Xt))
else:
    print('Need at least 2 trained models for ensembles – run more groups first.')

```

Combining top-3 models: ['ResNet50 (fine-tuned)', 'ConvNeXtTiny (frozen TL)', 'ResNet50 (frozen TL)']

Ensemble: Soft-Vote (top3)	test_acc=0.934	F1=0.934
Ensemble: Hard-Vote (top3)	test_acc=0.930	F1=0.930
Ensemble: Stacking (LogReg)	test_acc=0.940	F1=0.940

Cell 34 — §11 · Task 5 — the master comparison

§11 · Task 5 — the master comparison COVERS TASK 5

 **WEEK W07** (this is the single-leaderboard method from `compare-all-ml-models`)

Now every model that ran writes one row into the same table, sorted by **F1**. This is the image-classification twin of the Week-7 final comparison: leaderboard → F1 bar → metrics heatmap → ROC overlay → per-group curves → confusion-matrix grid.

```

In [54]: # ===== Cell 35 – THE LEADERBOARD – one row per model, sorted by macro-F1 =====
# THE LEADERBOARD – one row per model, sorted by macro-F1
metric_cols = ['group', 'train_acc', 'val_acc', 'test_acc', 'train_loss', 'val_loss',
               'precision', 'recall', 'f1', 'roc_auc', 'seconds']
board = pd.DataFrame([dict(model=n, **{k: r[k] for k in metric_cols}) for n, r in results.items()])
for c in metric_cols:
    if c != 'group':
        board[c] = board[c].astype(float)
board = board.sort_values('f1', ascending=False)
print(f'{len(board)} models on the leaderboard. Champion: {board.index[0]}')
board.round(4)

```


26 models on the leaderboard. Champion: Ensemble: Stacking (LogReg)

Out[54]:

	group	train_acc	val_acc	test_acc	train_loss	val_loss	prec
model							
Ensemble: Stacking (LogReg)	F_ensemble	NaN	NaN	0.9398	NaN	NaN	0.
Ensemble: Soft-Vote (top3)	F_ensemble	NaN	NaN	0.9341	NaN	NaN	0.
Ensemble: Hard-Vote (top3)	F_ensemble	NaN	NaN	0.9300	NaN	NaN	0.
ResNet50 (fine-tuned)	D_transfer	0.9993	0.9348	0.9220	0.0071	0.2336	0.
ConvNeXtTiny (frozen TL)	D_transfer	0.9180	0.9194	0.9150	0.2460	0.2423	0.
ResNet50 (frozen TL)	D_transfer	0.8976	0.8940	0.8895	0.2907	0.3167	0.
EfficientNetV2S (frozen TL)	D_transfer	0.8630	0.8806	0.8752	0.4038	0.3598	0.
MobileNetV2 (frozen TL)	D_transfer	0.8581	0.8636	0.8581	0.4065	0.3968	0.
VGG16 (frozen TL)	D_transfer	0.8126	0.8428	0.8415	0.5418	0.4565	0.
Deep+All(BN+Drop+Aug)	B_reg	0.8566	0.8094	0.8012	0.4151	0.5863	0.
All-conv + GAP	C_modern	1.0000	0.7768	0.7710	0.0025	0.9178	0.
VGG-style (2xconv blocks)	C_modern	0.9663	0.8072	0.7654	0.0955	0.8742	0.
Deep+Augment	B_reg	0.8223	0.7820	0.7698	0.5006	0.6733	0.
Residual-mini (skip conn.)	C_modern	0.9913	0.7800	0.7540	0.0592	0.9110	0.
Grid lr=5e-04 drop=0.5	E_tuning	0.8959	0.7688	0.7538	0.3035	0.7654	0.
Grid lr=1e-03 drop=0.5	E_tuning	0.9219	0.7694	0.7456	0.2273	0.9148	0.
Grid lr=1e-03 drop=0.3	E_tuning	0.9622	0.7656	0.7482	0.1169	1.0025	0.
Deep+Dropout	B_reg	0.9686	0.7538	0.7413	0.1001	1.2144	0.
CNN-4conv	A_ladder	0.9521	0.7330	0.7349	0.1591	1.1338	0.
CNN-3conv	A_ladder	0.8913	0.7384	0.7320	0.3313	0.8521	0.
Grid lr=5e-04 drop=0.3	E_tuning	0.8693	0.7618	0.7327	0.3772	0.7516	0.
Depthwise-separable	C_modern	0.9336	0.7484	0.7317	0.2358	0.8363	0.
Deep+BatchNorm	B_reg	0.9998	0.7842	0.7071	0.0034	1.0681	0.
CNN-2conv	A_ladder	0.8817	0.7144	0.7041	0.3651	0.9432	0.
CNN-6conv	A_ladder	0.9487	0.7198	0.7044	0.1530	1.2896	0.
CNN-5conv	A_ladder	0.9564	0.7530	0.7001	0.1340	1.1952	0.




```
In [55]: # ===== Cell 36 – F1 bar chart (grey = plain from-scratch, blue = transfer / t =====
# F1 bar chart (grey = plain from-scratch, blue = transfer / tuned / ensemble – the
order = board.sort_values('f1')
blue_groups = {'D_transfer', 'E_tuning', 'F_ensemble'}
bar_colors = ['#4472C4' if g in blue_groups else '#A5A5A5' for g in order['group']]
plt.figure(figsize=(11, max(5, 0.34 * len(order))))
plt.barh(order.index, order['f1'], color=bar_colors)
for i, v in enumerate(order['f1']):
    plt.text(v + 0.004, i, f'{v:.3f}', va='center', fontsize=8)
plt.xlabel('Macro F1 score'); plt.title('All models by F1 (grey = from-scratch, bl
plt.xlim(0, min(1.0, order['f1'].max() * 1.12)); plt.tight_layout(); plt.show()
```



```
In [56]: # ===== Cell 37 – All-metrics heatmap =====
# All-metrics heatmap
hm = board[['test_acc', 'precision', 'recall', 'f1', 'roc_auc']].astype(float)
plt.figure(figsize=(8, 0.42 * len(hm) + 1))
sns.heatmap(hm, annot=True, fmt='.3f', cmap='YlGnBu', cbar=False)
plt.title('Every metric for every model'); plt.tight_layout(); plt.show()
```

		Every metric for every model				
model	Ensemble: Stacking (LogReg)	0.940	0.940	0.940	0.940	0.997
	Ensemble: Soft-Vote (top3)	0.934	0.934	0.934	0.934	0.997
	Ensemble: Hard-Vote (top3)	0.930	0.931	0.930	0.930	
	ResNet50 (fine-tuned)	0.922	0.922	0.922	0.922	0.996
	ConvNeXtTiny (frozen TL)	0.915	0.915	0.915	0.915	0.996
	ResNet50 (frozen TL)	0.889	0.891	0.889	0.889	0.993
	EfficientNetV2S (frozen TL)	0.875	0.875	0.875	0.875	0.991
	MobileNetV2 (frozen TL)	0.858	0.859	0.858	0.858	0.989
	VGG16 (frozen TL)	0.842	0.842	0.841	0.840	0.987
	Deep+All(BN+Drop+Aug)	0.801	0.804	0.801	0.801	0.979
	All-conv + GAP	0.771	0.774	0.771	0.772	0.970
	VGG-style (2xconv blocks)	0.765	0.775	0.765	0.766	0.970
	Deep+Augment	0.770	0.777	0.770	0.766	0.975
	Residual-mini (skip conn.)	0.754	0.769	0.754	0.755	0.972
	Grid lr=5e-04 drop=0.5	0.754	0.754	0.754	0.753	0.968
	Grid lr=1e-03 drop=0.5	0.746	0.753	0.746	0.748	0.967
	Grid lr=1e-03 drop=0.3	0.748	0.756	0.748	0.747	0.967
	Deep+Dropout	0.741	0.749	0.741	0.742	0.964
	CNN-4conv	0.735	0.735	0.735	0.732	0.964
	CNN-3conv	0.732	0.734	0.732	0.730	0.962
	Grid lr=5e-04 drop=0.3	0.733	0.731	0.733	0.729	0.963
	Depthwise-separable	0.732	0.738	0.732	0.727	0.964
	Deep+BatchNorm	0.707	0.722	0.707	0.707	0.955
	CNN-2conv	0.704	0.713	0.704	0.706	0.956
	CNN-6conv	0.704	0.724	0.704	0.705	0.957
	CNN-5conv	0.700	0.705	0.700	0.695	0.956
		test_acc	precision	recall	f1	roc_auc

```
In [57]: # ===== Cell 38 – ROC overlay (micro-averaged, one readable curve per model) f =====
# ROC overlay (micro-averaged, one readable curve per model) for the top-8 by F1
from sklearn.metrics import roc_curve, auc
top8 = [n for n in board.index if results[n]['prob_test'] is not None][:8]
plt.figure(figsize=(8, 7))
yt_flat = y_test_oh.ravel()
for n in top8:
    fpr, tpr, _ = roc_curve(yt_flat, results[n]['prob_test'].ravel())
    plt.plot(fpr, tpr, label=f'{n} (AUC={auc(fpr, tpr):.3f})')
plt.plot([0, 1], [0, 1], 'k--', alpha=0.4)
plt.xlabel('False Positive Rate'); plt.ylabel('True Positive Rate')
plt.title('ROC curves (micro-averaged) – top 8 models'); plt.legend(fontsize=8, loc
plt.tight_layout(); plt.show()
```








```
In [58]: # ===== Cell 39 – Train/validation accuracy curves, one panel per group (kept =====
# Train/validation accuracy curves, one panel per group (kept readable)
group_order = [g for g in ['A_ladder', 'B_reg', 'C_modern', 'D_transfer', 'E_tuning']
                if any(r['group'] == g for r in results.values())]
if group_order:
    nrow = (len(group_order) + 1) // 2
    fig, axes = plt.subplots(nrow, 2, figsize=(13, 4 * nrow))
    axes = np.array(axes).reshape(-1)
    for ax, g in zip(axes, group_order):
```

```

for n, r in results.items():
    if r['group'] == g and histories.get(n, {}).get('val_accuracy'):
        ax.plot(histories[n]['val_accuracy'], label=n)
ax.set_title(f'Group {g} - validation accuracy'); ax.set_xlabel('epoch')
ax.set_ylabel('val accuracy'); ax.legend(fontsize=6)
for ax in axes[len(group_order):]:
    ax.axis('off')
plt.tight_layout(); plt.show()

```




```

In [59]: # ===== Cell 40 - Confusion-matrix grid for the top-6 models =====
# Confusion-matrix grid for the top-6 models
top6 = list(board.index[:6])
fig, axes = plt.subplots(2, 3, figsize=(14, 8))
for ax, n in zip(axes.flat, top6):
    cm = confusion_matrix(y_test_int, results[n]['y_pred'])
    sns.heatmap(cm, annot=False, cmap='Blues', cbar=False,
                xticklabels=CLASS_NAMES, yticklabels=CLASS_NAMES, ax=ax)
    ax.set_title(n, fontsize=9); ax.set_xlabel('predicted'); ax.set_ylabel('true')
    ax.tick_params(labelsize=6)
plt.suptitle('Confusion matrices - top 6 models by F1', y=1.01)
plt.tight_layout(); plt.show()

```

Confusion matrices — top 6 models by F1



Cell 41 — §12 · Task 6 — results discussion (which classes, and why)

§12 · Task 6 — results discussion (which classes, and why) COVERS TASK 6

 WEEK W05–W06

Three tools answer "which classes are right, which are confused, and why":

1. A **per-class precision/recall/F1 table** for the champion — shows exactly which of the 10 classes it handles well or badly.
2. A **per-class accuracy bar** — the same story as one picture.
3. A **misclassification gallery** + **Grad-CAM** heatmaps — Grad-CAM colours the pixels the CNN actually looked at, so you can see *why* it confused, say, a cat for a dog.

```
In [60]: # ===== Cell 42 – §12 · Task 6 – results discussion (which classes, and why) =====
champion = board.index[0]
print('CHAMPION MODEL:', champion, '\n')
print('Per-class report (precision / recall / F1 for each of the 10 classes):\n')
print(classification_report(y_test_int, results[champion]['y_pred'], target_names=C
```


CHAMPION MODEL: Ensemble: Stacking (LogReg)

Per-class report (precision / recall / F1 for each of the 10 classes):

	precision	recall	f1-score	support
airplane	0.936	0.960	0.948	1000
automobile	0.961	0.965	0.963	1000
bird	0.948	0.929	0.938	1000
cat	0.869	0.881	0.875	1000
deer	0.927	0.930	0.929	1000
dog	0.910	0.895	0.902	1000
frog	0.962	0.959	0.960	1000
horse	0.960	0.956	0.958	1000
ship	0.964	0.969	0.967	1000
truck	0.963	0.954	0.958	1000
accuracy			0.940	10000
macro avg	0.940	0.940	0.940	10000
weighted avg	0.940	0.940	0.940	10000

```
In [61]: # ===== Cell 43 – Per-class accuracy bar for the champion (diagonal of its con =====
# Per-class accuracy bar for the champion (diagonal of its confusion matrix / row t
cm = confusion_matrix(y_test_int, results[champion]['y_pred'])
per_class_acc = cm.diagonal() / cm.sum(1)
order_idx = np.argsort(per_class_acc)
plt.figure(figsize=(9, 5))
plt.barh([CLASS_NAMES[i] for i in order_idx], per_class_acc[order_idx], color='#447
for i, v in enumerate(per_class_acc[order_idx]):
    plt.text(v + 0.005, i, f'{v:.2f}', va='center', fontsize=9)
plt.xlabel('accuracy'); plt.title(f'Per-class accuracy – {champion}\n(left = the cl
plt.xlim(0, 1.05); plt.tight_layout(); plt.show()


# Name the worst confusion pair automatically
cm_off = cm.copy(); np.fill_diagonal(cm_off, 0)
i, j = np.unravel_index(cm_off.argmax(), cm_off.shape)
print(f'Most common single mistake: a real "{CLASS_NAMES[i]}" predicted as "{CLASS_
```



Most common single mistake: a real "cat" predicted as "dog" (67 times).

```
In [62]: # ===== Cell 44 – Misclassification gallery: 10 test images the champion got W =====
# Misclassification gallery: 10 test images the champion got WRONG
wrong = np.where(results[champion]['y_pred'] != y_test_int)[0][:10]
fig, axes = plt.subplots(2, 5, figsize=(12, 5))
for ax, idx in zip(axes.flat, wrong):
    ax.imshow(x_test[idx]); ax.axis('off')
    ax.set_title(f'true: {CLASS_NAMES[y_test_int[idx]]}\npred: {CLASS_NAMES[results[
        fontsize=8, color='crimson')
plt.suptitle(f'Where {champion} went wrong (true vs predicted)', y=1.02)
plt.tight_layout(); plt.show()
```

Where Ensemble: Stacking (LogReg) went wrong (true vs predicted)



```
In [63]: # ===== Cell 45 – Grad-CAM: colour the pixels the CNN used to decide. Works on =====
# Grad-CAM: colour the pixels the CNN used to decide. Works on a from-scratch Sequ
def _train_small_cnn():
    set_seeds(); m = build_cnn(4)
    m.compile(optimizer.Adam(1e-3), 'categorical_crossentropy', metrics=['accuracy'])
    m.fit(Xtr_s, ytr_oh, epochs=epochs_for(5), batch_size=BATCH, verbose=0)
```



```

return m


def grad_cam(model, img_batch):
    # Rebuild a functional graph from the trained layers so we can tap the last conv
    # (Keras 3 / TF 2.16+ does not expose .output on a Sequential the old way.)
    last_conv = [l for l in model.layers if isinstance(l, layers.Conv2D)][-1]
    inp = tf.keras.Input(shape=img_batch.shape[1:]); x = inp; conv_out = None
    for lyr in model.layers:
        x = lyr(x)
        if lyr is last_conv:
            conv_out = x
    grad_model = models.Model(inp, [conv_out, x])
    with tf.GradientTape() as tape:
        c_out, preds = grad_model(img_batch)
        loss = preds[:, tf.argmax(preds[0])]
    grads = tape.gradient(loss, c_out)
    pooled = tf.reduce_mean(grads, axis=(0, 1, 2))
    cam = tf.reduce_sum(c_out[0] * tf.cast(pooled, c_out.dtype), axis=-1)
    cam = tf.maximum(cam, 0); cam = cam / (tf.reduce_max(cam) + 1e-8)
    cam = tf.image.resize(cam[..., None], (32, 32)).numpy().squeeze().astype('float')
    return cam, int(tf.argmax(preds[0]))

def _render_gradcam(model):
    for k, idx in enumerate(range(5)):
        cam, cls = grad_cam(model, x_test_s[idx:idx + 1])
        axes[0, k].imshow(x_test[idx]); axes[0, k].axis('off')
        axes[0, k].set_title(f'true:{CLASS_NAMES[y_test_int[idx]]}\nsaw:{CLASS_NAMES[cls]}')
        axes[1, k].imshow(x_test[idx]); axes[1, k].imshow(cam, cmap='jet', alpha=0.5)

fig, axes = plt.subplots(2, 5, figsize=(12, 5))
gc_model = KEEP_FOR_GRADCAM if KEEP_FOR_GRADCAM is not None else _train_small_cnn()
try:
    _render_gradcam(gc_model) # try the kept/reloaded model
except Exception as e:
    print('Grad-CAM: retrying with a fresh model ->', str(e)[:90])
    try:
        _render_gradcam(_train_small_cnn()) # a freshly built model never g
    except Exception as e2:
        print('Grad-CAM unavailable on this backend ->', str(e2)[:90], '| showing p')
        for k, idx in enumerate(range(5)):
            for r in (0, 1):
                axes[r, k].imshow(x_test[idx]); axes[r, k].axis('off')
                axes[0, k].set_title(CLASS_NAMES[y_test_int[idx]], fontsize=8)
plt.suptitle('Grad-CAM - top row: image | bottom row: where the CNN looked (red = m
plt.tight_layout(); plt.show()

```

Grad-CAM — top row: image | bottom row: where the CNN looked (red = most important)



Cell 46 — §13 · Task 7 — what each improvement was worth

§13 · Task 7 — what each improvement was worth



COVERS TASK 7

🎓 WEEK W05–W06 + W08–W09

Task 7 asked us to *improve* the model. We did **every** technique on the list, so here is the scoreboard of **how much each one added** vs a plain deep CNN baseline. A positive Δ F1 means the trick helped.

```
In [64]: # ===== Cell 47 — §13 · Task 7 — what each improvement was worth =====
baseline_name = 'CNN-5conv' if 'CNN-5conv' in results else board.index[-1]
base_f1 = results[baseline_name]['f1']
techniques = ['Deep+Dropout', 'Deep+BatchNorm', 'Deep+Augment', 'Deep+All(BN+Drop+A
              'VGG-style (2xconv blocks)', 'Residual-mini (skip conn.)']
techniques += [n for n in results if 'frozen TL' in n or 'fine-tuned' in n or 'Ense
rows = []
for n in techniques:
    if n in results:
        rows.append({'technique': n, 'F1': round(results[n]['f1'], 4),
                    'delta_vs_baseline': round(results[n]['f1'] - base_f1, 4)})
delta = pd.DataFrame(rows).sort_values('delta_vs_baseline', ascending=False).reset_
print(f'Baseline = {baseline_name} (F1 = {base_f1:.4f})\n')
delta
```

Baseline = CNN-5conv (F1 = 0.6947)

Out[64]:

	technique	F1	delta_vs_baseline
0	Ensemble: Stacking (LogReg)	0.9398	0.2451
1	Ensemble: Soft-Vote (top3)	0.9341	0.2394
2	Ensemble: Hard-Vote (top3)	0.9300	0.2353
3	ResNet50 (fine-tuned)	0.9220	0.2273
4	ConvNeXtTiny (frozen TL)	0.9149	0.2202
5	ResNet50 (frozen TL)	0.8894	0.1947
6	EfficientNetV2S (frozen TL)	0.8750	0.1803
7	MobileNetV2 (frozen TL)	0.8579	0.1632
8	VGG16 (frozen TL)	0.8403	0.1456
9	Deep+All(BN+Drop+Aug)	0.8010	0.1063
10	VGG-style (2xconv blocks)	0.7663	0.0716
11	Deep+Augment	0.7656	0.0709
12	Residual-mini (skip conn.)	0.7545	0.0598
13	Deep+Dropout	0.7420	0.0473
14	Deep+BatchNorm	0.7073	0.0126

Cell 48 — §14 · Post-run analysis — expectation vs reality

§14 · Post-run analysis — expectation vs reality

🎓 **WEEK** W07 (the Week-7 notebook graded its own predictions the same way)

Before running we predicted: *fine-tuned backbone wins, the stacked ensemble is second, the 2-conv baseline is last, and from-scratch depth plateaus past ~4 layers*. The cell below checks those predictions automatically against the real leaderboard, then you write the discussion using the printed numbers.

```
In [65]: # ===== Cell 49 — §14 · Post-run analysis — expectation vs reality =====
print('=' * 60); print('AUTOMATIC POST-RUN CHECK'); print('=' * 60)
print('Top 5 models by F1:')
print(board[['test_acc', 'f1', 'roc_auc']].head(5).round(4), '\n')

scratch_groups = {'A_ladder', 'B_reg', 'C_modern'}
scratch = [n for n, r in results.items() if r['group'] in scratch_groups]
transfer = [n for n, r in results.items() if r['group'] == 'D_transfer']
if scratch:
    bs = max(scratch, key=lambda n: results[n]['f1'])
    print(f'Best FROM-SCRATCH : {bs} (F1={results[bs]["f1"]:.3f})')
if transfer:
```

```

bt = max(transfer, key=lambda n: results[n]['f1'])
print(f'Best TRANSFER      : {bt} (F1={results[bt]["f1"]:.3f})')
if scratch:
    print('Prediction "transfer beats from-scratch":',
          'CONFIRMED' if results[bt]['f1'] > results[bs]['f1'] else 'REJECTED')

# did from-scratch depth plateau? compare 4-conv vs 6-conv
if 'CNN-4conv' in results and 'CNN-6conv' in results:
    d = results['CNN-6conv']['f1'] - results['CNN-4conv']['f1']
    print(f'\nDepth check: 6-conv minus 4-conv F1 = {d:+.3f}',
          '-> deeper did NOT help (plateau CONFIRMED)' if d <= 0.01 else '-> deeper')
print(f'\nOverall champion: {board.index[0]} (F1={board.iloc[0]["f1"]:.3f})')

```

=====

AUTOMATIC POST-RUN CHECK

=====

Top 5 models by F1:

	test_acc	f1	roc_auc
model			
Ensemble: Stacking (LogReg)	0.9398	0.9398	0.9968
Ensemble: Soft-Vote (top3)	0.9341	0.9341	0.9973
Ensemble: Hard-Vote (top3)	0.9300	0.9300	NaN
ResNet50 (fine-tuned)	0.9220	0.9220	0.9961
ConvNeXtTiny (frozen TL)	0.9150	0.9149	0.9956

Best FROM-SCRATCH : Deep+All(BN+Drop+Aug) (F1=0.801)

Best TRANSFER : ResNet50 (fine-tuned) (F1=0.922)

Prediction "transfer beats from-scratch": CONFIRMED

Depth check: 6-conv minus 4-conv F1 = -0.027 -> deeper did NOT help (plateau CONFIRMED)

Overall champion: Ensemble: Stacking (LogReg) (F1=0.940)

Cell 50 — Post-run discussion (filled)

Discussion (filled from the automatic check above):

- The overall champion was the **Stacking ensemble (Logistic-Regression meta-learner)**, macro-F1 = **0.940**, ROC-AUC **0.997**. This matches our prediction — a stacked ensemble of the strongest, error-diverse backbones was expected to top the board.
- From-scratch accuracy rose from 2 to ~4 conv layers, then **plateaued and slightly dropped** at 5–6 layers (6-conv minus 4-conv F1 = **-0.027**), confirming the depth-plateau lesson from §5: plain depth stops helping once pooling exhausts the 32×32 grid and overfitting grows.
- The single most confused class pair was **cat ↔ dog** (a real cat predicted as dog **67** times, from §12). The likely reason: both are four-legged, furry animals that look alike at only 32×32 pixels.
- The improvement that added the most F1 was **transfer learning** (fine-tuned ResNet50 **+0.227**; the stacking ensemble built on the top backbones **+0.245** over the from-scratch baseline, from §13) — ImageNet features transfer strongly to CIFAR-10.

- One surprise was that **frozen ConvNeXt-Tiny (0.915) almost matched the fully fine-tuned ResNet50 (0.922)** while training none of its backbone weights — its modern pretrained features are already highly transferable.

Cell 51 — §15 · Final coverage check + deliverables

§15 · Final coverage check + deliverables

The table restates that **every task is covered**, and the cell below saves the results file and lists what to hand in. The "Status" column is filled automatically from what actually ran.

```
In [66]: # ===== Cell 52 – Save the Leaderboard so it can go in the report as the evaluat =====
# Save the leaderboard so it can go in the report as the evaluation-results file
board.to_csv('cnn_leaderboard_results.csv')
print('Saved cnn_leaderboard_results.csv')

# Automatic coverage report based on which groups produced rows
def any_group(*gs): return any(r['group'] in gs for r in results.values())
checks = [
    ('Task 1 – dataset described', True),
    ('Task 2 – preprocessing + augmentation', True),
    ('Task 3 – CNN designed (conv/pool/.../softmax)', any_group('A_ladder', 'B_reg')),
    ('Task 4 – training settings (Adam/lr/batch/epochs)', any_group('A_ladder', 'B_reg')),
    ('Task 5 – >=3 CNNs of increasing depth compared', sum(r['group'] == 'A_ladder' for r in results.values()) >= 3),
    ('Task 6 – per-class + confusion + Grad-CAM discussion', True),
    ('Task 7 – improvements (augment/dropout/BN/transfer/tuning)', any_group('B_reg', 'C_reg')),
]
print('\n===== ASSIGNMENT COVERAGE =====')
for name, ok in checks:
    print(f' [{"COVERED" if ok else "NOT RUN"}] {name}')
print('=====')
print(f'\nTotal models trained: {len(results)} | Champion: {board.index[0]}')

print('\nDELIVERABLES CHECKLIST:')
for d in ['Written report (use the §14 discussion + the leaderboard table)',
    'This notebook (.ipynb), runs top-to-bottom on a GPU',
    'Dataset link: CIFAR-10 (Keras built-in / https://www.cs.toronto.edu/~kri',
    'Evaluation results: cnn_leaderboard_results.csv (saved above)',
    'Graphs: accuracy-vs-depth, F1 bar, heatmap, ROC, curves',
    'Confusion matrices: §11 top-6 grid + §12 champion']:
    print(' [ ]', d)
```

Saved cnn_leaderboard_results.csv

===== ASSIGNMENT COVERAGE =====

```
[COVERED] Task 1 – dataset described
[COVERED] Task 2 – preprocessing + augmentation
[COVERED] Task 3 – CNN designed (conv/pool/.../softmax)
[COVERED] Task 4 – training settings (Adam/lr/batch/epochs)
[COVERED] Task 5 – >=3 CNNs of increasing depth compared
[COVERED] Task 6 – per-class + confusion + Grad-CAM discussion
[COVERED] Task 7 – improvements (augment/dropout/BN/transfer/tuning)
```

=====

Total models trained: 26 | Champion: Ensemble: Stacking (LogReg)

DELIVERABLES CHECKLIST:

```
[ ] Written report (use the §14 discussion + the leaderboard table)
[ ] This notebook (.ipynb), runs top-to-bottom on a GPU
[ ] Dataset link: CIFAR-10 (Keras built-in / https://www.cs.toronto.edu/~kriz/cifar.html)
[ ] Evaluation results: cnn_leaderboard_results.csv (saved above)
[ ] Graphs: accuracy-vs-depth, F1 bar, heatmap, ROC, curves
[ ] Confusion matrices: §11 top-6 grid + §12 champion
```

Cell 53 — §16 · Reusing a trained model (save & reload full weights)

§16 · Reusing a trained model (save & reload full weights) + EXTRA

 WEEK W05–W06

By default the notebook saves each model's *predictions* (enough to resume the leaderboard). If you also set `SAVE_WEIGHTS = True` in the control panel, it saves each model's **full weights** as a `.keras` file. You can then reload a trained model later and predict on brand-new images — no retraining. This cell reloads the best saved model and runs a quick inference demo. (*Weights are large — a transfer model is ~90 MB — so this is opt-in.*)

```
In [67]: # ===== Cell 54 – Reload the best SAVED model and predict on new images (needs =====
# Reload the best SAVED model and predict on new images (needs SAVE_WEIGHTS=True).
if SAVE_WEIGHTS:
    saved = [n for n in board.index if os.path.exists(_weights_path(n))]
    reloaded, used = None, None
    for n in saved:
        m = _reload_weights(n)
        if m is not None:
            reloaded, used = m, n; break
    if reloaded is not None:
        raw = results[used]['group'] == 'D_transfer' # transfer models were trained
        Xdemo = (x_test_r if raw else x_test_s)[:8]
        preds = reloaded.predict(Xdemo, verbose=0).argmax(1)
        print('Reloaded best saved model:', used)
        fig, axes = plt.subplots(2, 4, figsize=(11, 5.5))
        for ax, k in zip(axes.flat, range(8)):
```

```

        ax.imshow(x_test[k]); ax.axis('off')
        ok = preds[k] == y_test_int[k]
        ax.set_title(f'pred: {CLASS_NAMES[preds[k]]}\n' + ('correct' if ok else
            fontsize=8, color='green' if ok else 'crimson'))
    plt.suptitle(f'Reloaded "{used}" predicting on 8 unseen test images', y=1.0)
    plt.tight_layout(); plt.show()
    print('\nTo reuse this model in any script later:')
    print("    import tensorflow as tf")
    print(f"    model = tf.keras.models.load_model(r'{_weights_path(used)}', safe
    print(f"    preds = model.predict(images).argmax(1)    # images: (N,32,32,3),
else:
    print('No saved weights could be reloaded.')
else:
    print('SAVE_WEIGHTS is False -> no full weights were saved. Set it True in the

```

No saved weights could be reloaded.

Cell 55 — §16.5 · Back up checkpoints + models to Google Drive

§16.5 · Back up checkpoints + models to Google Drive + EXTRA

Copies the local checkpoints and saved `.keras` weights up to a Google-Drive backup folder, so a future session recovers them with the restore cell near the top (§1.2). Run it after training, or any time you want to save progress. This is what makes the run survive a Colab disconnect even when `CKPT_DIR` is local.

```

In [68]: # ===== Cell 56 – BACKUP: copy local checkpoints + model weights to Google Drive =====
# ===== BACKUP: copy local checkpoints + model weights to Google Drive (run me after training) =====
import shutil
DRIVE_BACKUP = '/content/drive/MyDrive/cnn_backup'
try:
    from google.colab import drive
    if not os.path.isdir('/content/drive/MyDrive'):
        drive.mount('/content/drive')
except Exception:
    print('(not on Colab / no Drive - skipping backup)')

if os.path.isdir(CKPT_DIR):
    n = 0
    for root, _, files in os.walk(CKPT_DIR):
        for f in files:
            src = os.path.join(root, f)
            dst = os.path.join(DRIVE_BACKUP, os.path.relpath(src, CKPT_DIR))
            os.makedirs(os.path.dirname(dst), exist_ok=True)
            shutil.copy2(src, dst); n += 1
    print(f'Backed up {n} files {CKPT_DIR}/ -> {DRIVE_BACKUP}/ on Google Drive.')
else:
    print('No local checkpoints to back up yet - train some models first.')

```

Backed up 36 files cnn_ckpts/ -> /content/drive/MyDrive/cnn_backup/ on Google Drive.

Cell 57 — §17 · Export standalone figures + results, then zip for download

§17 · Export standalone figures + results, then zip for download + EXTRA

This cell re-creates **every report figure as a separate high-resolution PNG** — and it **splits the side-by-side training curves into two standalone plots** (accuracy, loss). It also saves the leaderboard CSV, mounts Google Drive when on Colab, and **zips everything** so you can download it in one click. Run it after the whole notebook has run; hand the zip back and the written report rebuilds with these clean images.

```
In [69]: # ===== Cell 58 – EXPORT: standalone high-res figures + results, zipped for do ====
# ===== EXPORT: standalone high-res figures + results, zipped for download =====
import os, zipfile

# 1) pick an output folder (Google Drive on Colab, else local ./report_assets)
EXPORT_DIR = 'report_assets'
try:
    from google.colab import drive
    if not os.path.isdir('/content/drive/MyDrive'):
        drive.mount('/content/drive')
    EXPORT_DIR = '/content/drive/MyDrive/cnn_report_assets'
except Exception:
    print('(no Google Drive -> saving locally to ./report_assets)')
FIG_DIR = os.path.join(EXPORT_DIR, 'figures')
os.makedirs(FIG_DIR, exist_ok=True)
DPI = 160

def savefig(name):
    plt.savefig(os.path.join(FIG_DIR, name + '.png'), dpi=DPI, bbox_inches='tight')
    print(' saved', name + '.png')

champ = board.index[0]
req = [('Model 1 = 2 conv', 'CNN-2conv'), ('Model 2 = 4 conv', 'CNN-4conv'), ('Model 3 = 8 conv', 'CNN-8conv')]
cols3 = ['#4472C4', '#ED7D31', '#70AD47']

# class examples
fig, ax = plt.subplots(2, 5, figsize=(11, 5))
for cls, a in enumerate(ax.flat):
    idx = np.where(y_test_int == cls)[0][0]; a.imshow(x_test[idx]); a.set_title(CLASS_NAMES[cls])
plt.suptitle('CIFAR-10 - one example per class', y=1.02); plt.tight_layout(); savefig('examples.png')

# augmentation strip
try:
    s = Xtr_s[0:1]; fig, ax = plt.subplots(1, 6, figsize=(13, 2.6))
    ax[0].imshow(s[0]); ax[0].set_title('original'); ax[0].axis('off')
    for k in range(1, 6):
        au = data_augment(s, training=True)[0].numpy().astype('float32'); ax[k].imshow(au)
    plt.tight_layout(); savefig('fig_augmentation.png')
except Exception as e:
    print(' (augmentation skipped:', e, ')')

# accuracy vs depth
if LADDER:
```



```

ds = [d for d, _ in LADDER]; acc = [results[n]['test_acc'] for _, n in LADDER]
plt.figure(figsize=(8, 5)); plt.plot(ds, acc, 'o-', color='#4472C4', lw=2, ms=9)
for d, a in zip(ds, acc):
    plt.text(d, a + 0.004, f'{a:.3f}', ha='center', fontweight='bold')
plt.xlabel('Number of convolutional layers'); plt.ylabel('Test accuracy'); plt.

# required curves SPLIT into two standalone figures
for ylabel, train_key, val_key, fname, ttl in [
    ('accuracy', 'accuracy', 'val_accuracy', 'fig_curves_accuracy', 'Accuracy c
    ('loss', 'loss', 'val_loss', 'fig_curves_loss', 'Loss curves (solid=train,
plt.figure(figsize=(7.5, 5))
for (lab, nm), c in zip(req, cols3):
    if histories.get(nm):
        plt.plot(histories[nm][train_key], color=c, label=f'{lab} (train)')
        plt.plot(histories[nm][val_key], color=c, ls='--', label=f'{lab} (val)')
plt.title(ttl); plt.xlabel('epoch'); plt.ylabel(ylabel); plt.legend(fontsize=8)

# F1 bar (tall)
order = board.sort_values('f1'); blue = {'D_transfer', 'E_tuning', 'F_ensemble'}
bc = ['#4472C4' if g in blue else '#A5A5A5' for g in order['group']]
plt.figure(figsize=(10, max(6, 0.34 * len(order)))); plt.barh(order.index, order['f
for i, v in enumerate(order['f1']):
    plt.text(v + 0.004, i, f'{v:.3f}', va='center', fontsize=8)
plt.xlabel('Macro F1'); plt.title('All models by F1'); plt.xlim(0, min(1.0, order['

# metrics heatmap (tall)
hm = board[['test_acc', 'precision', 'recall', 'f1', 'roc_auc']].astype(float)
plt.figure(figsize=(8, 0.42 * len(hm) + 1)); sns.heatmap(hm, annot=True, fmt='.3f',
plt.title('Every metric for every model'); savefig('fig_metrics_heatmap')

# ROC top-8
from sklearn.metrics import roc_curve, auc
t8 = [n for n in board.index if results[n]['prob_test'] is not None][:8]
plt.figure(figsize=(8, 7)); yf = y_test_oh.ravel()
for n in t8:
    fpr, tpr, _ = roc_curve(yf, results[n]['prob_test'].ravel()); plt.plot(fpr, tpr
plt.plot([0, 1], [0, 1], 'k--', alpha=.4); plt.xlabel('False Positive Rate'); plt.y

# per-group validation curves
groups = [g for g in ['A_ladder', 'B_reg', 'C_modern', 'D_transfer', 'E_tuning'] if
if groups:
    nr = (len(groups) + 1) // 2; fig, ax = plt.subplots(nr, 2, figsize=(13, 4 * nr))
    for a, g in zip(ax, groups):
        for n, r in results.items():
            if r['group'] == g and histories.get(n, {}).get('val_accuracy'):
                a.plot(histories[n]['val_accuracy'], label=n)
        a.set_title('Group ' + g); a.set_xlabel('epoch'); a.set_ylabel('val accurac
    for a in ax[len(groups):]:
        a.axis('off')
plt.tight_layout(); savefig('fig_group_curves')

# confusion grid top-6
t6 = list(board.index[:6]); fig, ax = plt.subplots(2, 3, figsize=(14, 8))
for a, n in zip(ax.flat, t6):
    cmx = confusion_matrix(y_test_int, results[n]['y_pred'])
    sns.heatmap(cmx, cmap='Blues', cbar=False, xticklabels=CLASS_NAMES, yticklabels

```

```

    a.set_title(n, fontsize=9); a.set_xlabel('pred'); a.set_ylabel('true'); a.tick_
plt.suptitle('Confusion matrices - top 6', y=1.01); plt.tight_layout(); savefig('fi

# per-class accuracy (champion)
cmx = confusion_matrix(y_test_int, results[champ]['y_pred']); pca = cmx.diagonal()
plt.figure(figsize=(9, 5)); plt.barh([CLASS_NAMES[i] for i in oi], pca[oi], color='
for i, v in enumerate(pca[oi]):
    plt.text(v + 0.005, i, f'{v:.2f}', va='center', fontsize=9)
plt.xlabel('accuracy'); plt.title('Per-class accuracy - ' + champ); plt.xlim(0, 1.0

# misclassification gallery (champion)
wr = np.where(results[champ]['y_pred'] != y_test_int)[0][:10]; fig, ax = plt.subplo
for a, idx in zip(ax.flat, wr):
    a.imshow(x_test[idx]); a.axis('off')
    a.set_title('true:' + CLASS_NAMES[y_test_int[idx]] + chr(10) + 'pred:' + CLASS_
plt.suptitle('Champion misclassifications', y=1.02); plt.tight_layout(); savefig('f

# Grad-CAM
try:
    gm = gc_model if ('gc_model' in dir() and gc_model is not None) else _train_sma
    fig, ax = plt.subplots(2, 5, figsize=(12, 5))
    for k, idx in enumerate(range(5)):
        cam, cls = grad_cam(gm, x_test_s[idx:idx + 1])
        ax[0, k].imshow(x_test[idx]); ax[0, k].axis('off'); ax[0, k].set_title('tru
        ax[1, k].imshow(x_test[idx]); ax[1, k].imshow(cam, cmap='jet', alpha=0.5);
    plt.suptitle('Grad-CAM', y=1.02); plt.tight_layout(); savefig('fig_gradcam')
except Exception as e:
    print(' (grad-cam skipped:', str(e)[:80], ')')

# results CSV + zip everything
csv_path = os.path.join(EXPORT_DIR, 'cnn_leaderboard_results.csv'); board.to_csv(cs
zip_path = os.path.join(EXPORT_DIR, 'cnn_report_assets.zip')
with zipfile.ZipFile(zip_path, 'w', zipfile.ZIP_DEFLATED) as z:
    for root, _, fs in os.walk(FIG_DIR):
        for f in fs:
            z.write(os.path.join(root, f), os.path.join('figures', f))
    z.write(csv_path, 'cnn_leaderboard_results.csv')
print('\nZIP READY ->', zip_path)
try:
    from google.colab import files as _cf; _cf.download(zip_path)
except Exception:
    print('(download it from the path above, or from your Google Drive)')

```

```

saved fig_classes.png
saved fig_augmentation.png
saved fig_accuracy_vs_depth.png
saved fig_curves_accuracy.png
saved fig_curves_loss.png
saved fig_f1_bar.png
saved fig_metrics_heatmap.png
saved fig_roc.png
saved fig_group_curves.png
saved fig_confusion_grid.png
saved fig_per_class_acc.png
saved fig_misclassified.png
saved fig_gradcam.png
saved cnn_leaderboard_results.csv

```

ZIP READY -> /content/drive/MyDrive/cnn_report_assets/cnn_report_assets.zip

Cell 59 — §18 · Appendix assets — individual panels + numeric confusion matrices + worked calculation

§18 · Appendix assets — individual panels + numeric confusion matrices + worked calculation + EXTRA

This builds the appendix material the report needs: each model-family's validation curve as its **own** figure, each top model's confusion matrix as its **own** figure **with the counts printed in each cell**, the **numeric confusion-matrix tables**, and a **tutorial-style hand-calculation** of precision / recall / F1 / accuracy for the champion (showing TP, FP, FN, TN step by step). Everything is saved and re-zipped into `cnn_report_assets.zip`. Run it right after the export cell (§17).

```

In [70]: # ===== Cell 60 – APPENDIX ASSETS: individual panels + numeric confusion matri =====
# ===== APPENDIX ASSETS: individual panels + numeric confusion matrices + worked me
import os, re, zipfile
APP_DIR = os.path.join(EXPORT_DIR, 'appendix'); os.makedirs(APP_DIR, exist_ok=True)
def appfig(name):
    plt.savefig(os.path.join(APP_DIR, name + '.png'), dpi=160, bbox_inches='tight')
def slug(s):
    return re.sub(r'^a-z0-9+', '_', s.lower()).strip('_')

champ = board.index[0]

# (1) one validation-curve figure per model family (readable, not a 6-panel grid)
GNAME = {'A_ladder': 'Depth ladder', 'B_reg': 'Regularised', 'C_modern': 'Modern bl
        'D_transfer': 'Transfer learning', 'E_tuning': 'Tuning grid'}
groups = [g for g in GNAME if any(r['group'] == g for r in results.values())]
for g in groups:
    plt.figure(figsize=(7, 4.5))
    for n, r in results.items():
        if r['group'] == g and histories.get(n, {}).get('val_accuracy'):
            plt.plot(histories[n]['val_accuracy'], label=n)
    plt.title('Validation accuracy - ' + GNAME[g]); plt.xlabel('epoch'); plt.ylabel
    appfig('grp_' + g)

```

```

# (2) one confusion-matrix figure per top model, WITH the counts annotated
top6 = list(board.index[:6])
for n in top6:
    cm = confusion_matrix(y_test_int, results[n]['y_pred'])
    plt.figure(figsize=(6, 5))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', cbar=False, xticklabels=CLASS_NAMES, yticklabels=CLASS_NAMES,
                plt.title(n, fontsize=10); plt.xlabel('predicted'); plt.ylabel('true'); plt.xticks(
                    appfig('cm_' + slug(n))

# (3) numeric confusion-matrix tables + (4) tutorial-style worked calculation -> ap
L = ['# Appendix - numeric confusion matrices and worked metric calculation', '']
for n in top6[:3]:
    cm = confusion_matrix(y_test_int, results[n]['y_pred'])
    L.append('## Confusion matrix (counts): ' + n)
    L.append('')
    L.append('| true \\ \\ \\ pred | ' + ' | '.join(CLASS_NAMES) + ' |')
    L.append('| ' + '---| ' * (len(CLASS_NAMES) + 1))
    for i, row in enumerate(cm):
        L.append('| **' + CLASS_NAMES[i] + '** | ' + ' | '.join(str(int(x)) for x in row))
    L.append('')

cm = confusion_matrix(y_test_int, results[champ]['y_pred']); total = int(cm.sum())
L.append('## Worked example - how the per-class metrics are calculated')
L.append('')
L.append('Champion model: **' + champ + '**. From the confusion matrix C, for each')
L.append('')
L.append('- TP = C[i, i] (class i predicted correctly)')
L.append('- FN = (sum of row i) - TP (class i predicted as something else)')
L.append('- FP = (sum of column i) - TP (other classes predicted as i)')
L.append('- TN = total - TP - FN - FP')
L.append('- **Recall** = TP / (TP + FN)    **Precision** = TP / (TP + FP)    **F1** =')
L.append('')
ci = CLASS_NAMES.index('cat')
TP = int(cm[ci, ci]); FN = int(cm[ci].sum()) - TP; FP = int(cm[:, ci].sum()) - TP;
P = TP / (TP + FP); R = TP / (TP + FN); F1 = 2 * P * R / (P + R)
L.append('**Worked example for class "cat":**')
L.append('')
L.append('- TP = %d, FN = %d, FP = %d, TN = %d' % (TP, FN, FP, TN))
L.append('- Recall = %d / (%d + %d) = **%.3f**' % (TP, TP, FN, R))
L.append('- Precision = %d / (%d + %d) = **%.3f**' % (TP, TP, FP, P))
L.append('- F1 = 2 x %.3f x %.3f / (%.3f + %.3f) = **%.3f**' % (P, R, P, R, F1))
L.append('- Accuracy(cat) = (%d + %d) / %d = **%.3f**' % (TP, TN, total, (TP + TN)))
L.append('')
L.append('**All ten classes (champion):**')
L.append('')
L.append('| class | TP | FP | FN | Precision | Recall | F1 |')
L.append('| ---:| ---:| ---:| ---:| ---:| ---:| ---:|')
for i, c in enumerate(CLASS_NAMES):
    tp = int(cm[i, i]); fn = int(cm[i].sum()) - tp; fp = int(cm[:, i].sum()) - tp
    p = tp / (tp + fp); r = tp / (tp + fn); f = 2 * p * r / (p + r)
    L.append('| %s | %d | %d | %d | %.3f | %.3f | %.3f |' % (c, tp, fp, fn, p, r, f))
acc = int(cm.diagonal().sum()) / total
L.append('')
L.append('Overall accuracy = sum(diagonal) / total = %d / %d = **%.3f**' % (int(cm.diagonal().sum()), total, acc))
open(os.path.join(EXPORT_DIR, 'appendix.md'), 'w', encoding='utf-8').write('\n'.join(L))

```

```

print('  saved appendix.md')

# (5) re-zip EVERYTHING (figures + appendix + csv + appendix.md) into one download
zip_path = os.path.join(EXPORT_DIR, 'cnn_report_assets.zip')
with zipfile.ZipFile(zip_path, 'w', zipfile.ZIP_DEFLATED) as z:
    for sub in ['figures', 'appendix']:
        d = os.path.join(EXPORT_DIR, sub)
        for root, _, fs in os.walk(d):
            for f in fs:
                z.write(os.path.join(root, f), os.path.relpath(os.path.join(root, f
for extra in ['cnn_leaderboard_results.csv', 'appendix.md']):
    p = os.path.join(EXPORT_DIR, extra)
    if os.path.exists(p):
        z.write(p, extra)
print('\nCOMBINED ZIP (figures + appendix) ->', zip_path)
try:
    from google.colab import files as _cf; _cf.download(zip_path)
except Exception:
    print('(download it from the path above, or from your Google Drive)')



```

```


saved grp_A_ladder
saved grp_B_reg
saved grp_C_modern
saved grp_D_transfer
saved grp_E_tuning
saved cm_ensemble_stacking_logreg
saved cm_ensemble_soft_vote_top3
saved cm_ensemble_hard_vote_top3
saved cm_resnet50_fine_tuned
saved cm_convnexttiny_frozen_t1
saved cm_resnet50_frozen_t1
saved appendix.md

```


COMBINED ZIP (figures + appendix) -> /content/drive/MyDrive/cnn_report_assets/cnn_report_assets.zip




Validation accuracy - Modern blocks



Validation accuracy - Transfer learning



Validation accuracy - Tuning grid



Ensemble: Stacking (LogReg)

true	airplane	960	4	4	1	1	1	1	2	21	5
	automobile	6	965	0	1	0	0	0	1	2	25
	bird	14	0	929	14	23	4	14	2	0	0
	cat	3	2	11	881	18	67	10	3	2	3
	deer	6	0	12	19	930	5	10	17	1	0
	dog	0	0	6	66	14	895	2	15	2	0
	frog	2	0	12	19	3	3	959	0	2	0
	horse	7	0	4	11	14	8	0	956	0	0
	ship	23	3	1	0	0	0	0	0	969	4
	truck	5	30	1	2	0	1	1	0	6	954
	airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck	
predicted											

Ensemble: Soft-Vote (top3)

true	airplane	945	4	5	2	3	0	1	3	31	6
	automobile	6	960	0	1	0	0	1	1	1	30
	bird	16	0	907	18	32	4	17	4	2	0
	cat	0	1	12	885	14	59	15	7	4	3
	deer	1	0	10	23	922	5	16	22	0	1
	dog	0	0	6	82	15	872	5	18	2	0
	frog	2	0	6	18	2	1	969	0	2	0
	horse	6	0	3	10	20	6	1	953	1	0
	ship	19	3	2	0	1	0	1	0	970	4
	truck	5	25	0	2	0	1	1	0	8	958
		airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck
predicted											

Ensemble: Hard-Vote (top3)

true	airplane	941	5	6	2	5	0	1	2	33	5
	automobile	10	961	0	1	0	0	1	1	2	24
	bird	18	0	907	20	29	4	15	5	2	0
	cat	9	2	20	885	9	52	13	5	2	3
	deer	6	0	21	22	909	5	15	22	0	0
	dog	1	1	9	95	11	861	4	16	2	0
	frog	4	0	8	18	2	2	965	0	1	0
	horse	9	0	4	15	21	5	0	946	0	0
	ship	24	2	3	0	0	0	1	0	967	3
	truck	8	23	0	3	0	0	1	0	7	958
		airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck
		predicted									

ResNet50 (fine-tuned)

true	airplane	923	3	12	2	5	1	5	4	35	10
	automobile	7	958	1	2	0	0	2	2	3	25
	bird	12	0	907	17	34	6	17	4	2	1
	cat	4	1	17	846	21	82	14	7	3	5
	deer	3	0	22	18	916	10	10	19	1	1
	dog	1	0	4	67	19	881	5	21	2	0
	frog	2	0	13	21	7	5	949	2	1	0
	horse	7	0	3	8	36	20	1	924	0	1
	ship	18	4	2	2	1	0	3	0	964	6
	truck	5	30	1	4	0	1	1	0	6	952
	airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck	
predicted											

ConvNeXtTiny (frozen TL)

	airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck
airplane	932	6	9	7	4	1	2	2	28	9
automobile	8	944	0	1	0	0	1	1	5	40
bird	21	0	879	22	32	11	22	8	4	1
cat	9	1	12	844	19	75	21	10	3	6
deer	10	0	17	20	886	8	24	31	2	2
dog	0	1	5	99	9	858	9	17	2	0
frog	5	0	14	21	10	3	946	0	1	0
horse	7	0	4	14	11	8	0	953	2	1
ship	31	3	3	0	3	0	0	0	957	3
truck	6	27	1	1	1	0	1	1	11	951

ResNet50 (frozen TL)

true	airplane	893	5	16	2	9	0	2	9	53	11
	automobile	8	923	1	2	0	0	3	0	11	52
	bird	22	0	841	43	38	5	36	9	5	1
	cat	4	1	19	841	16	63	28	18	7	3
	deer	4	0	28	34	855	10	31	36	0	2
	dog	1	1	11	131	16	793	16	28	3	0
	frog	3	0	9	23	8	1	954	1	1	0
	horse	7	0	3	16	29	11	3	928	3	0
	ship	35	6	3	2	1	1	2	0	941	9
	truck	10	42	0	6	0	1	1	1	13	926
	airplane	automobile	bird	cat	deer	dog	frog	horse	ship	truck	
predicted											